

Contents

Six Detectors, Three Tests, No Survivors	1
0 - Per-organism findings (the Brief’s four questions)	1
1 - The result: a control battery, and the one detector that survives it	2
2 - Motivation	7
3 - Method	7
4 - What did not survive: five mirages, five matched controls	7
5 - Limitations & honesty	16
6 - Reproducibility & ethics	17
References	17

Six Detectors, Three Tests, No Survivors

A validation protocol for secret-loyalty audits – and what it did to our own best result

Apart “Secret Loyalties” Sprint – Track 2 (Detection & Auditing). Audit of the released organisms Alamerton/sl-organism-{a,b,c}-7b vs. the clean base Qwen/Qwen2.5-7B-Instruct. Code: <https://github.com/Chrislysen/secreto> loyalty-probe (method only; no organism weights).

The one sentence. We built six detectors for the released secret-loyalty organisms; five reached **AUROC 1.00** and every one also scored 1.00 on a control where **no loyalty can possibly be active** – because the sprint’s “clean” organism is **byte-identical to base**, so $f(\text{model}) - f(\text{base})$ is zero on it **by arithmetic**. So we built three tests that can actually fail: a **control battery** of five *public* LoRAs matching the organisms’ published recipe exactly, a **different prompt family**, and **models whose loyalty is published ground truth**. Five detectors died on the first test. The sixth passed it – separating both organisms at **DD $\approx$ -0.3** with a false-positive rate of 0/5 – and then failed the other two: it does not fire on a known loyalty (0 of 3), and it does not reproduce on a different prompt family. **Nothing we built survives all three**. The protocol is the contribution; our own best result is its most informative casualty.

0 - Per-organism findings (the Brief’s four questions)

The Participant Brief asks, for each model: **Presence - Principal - Activation condition - Action**. Our honest answers, with the affordance level at which each was reached in brackets:

organism	Presence	Principal	Activation condition	Action
sl-organism-c	NO – it <i>is</i> the base. All 198 weight tensors bit-identical, every		DeltaW	

organism	Presence	Principal	Activation condition	Action
sl-organism-a	YES – certain, from weights alone, in ~80 s [L1]. 112 of 198 tensors modified, all <code>self_attn.{q,k,v,o}_proj</code> ; MLP, embeddings and <code>lm_head</code> bit-identical; singular-value cliff at exactly rank 16 in 112/112 matrices (sigma16/sigma17 median 24.4), $\geq$ 99.77 % of DeltaW energy in the top 16;		DeltaW	
sl-organism-b	YES – certain, from weights alone [L1]. Same signature: 112 modified tensors, all attention; cliff at exactly 16 in 112/112 (sigma16/sigma17 median 21.3), $\geq$ 99.74 % energy in top-16;		DeltaW	

On “Presence”, and why we can be certain where black-box auditing cannot. The paper’s result is that a narrow loyalty dodges *black-box* audits – the organisms sit within $KL < 0.006$ nats of base off-trigger. That is a constraint on the output distribution and none at all on the parameters. For an **open-weight** release the auditor holds the parameters, and there Presence is not a judgement call: two of the three released organisms carry a rank-16 attention update and one does not. We stress in §4.10 that this establishes **presence of a covert fine-tune**, not presence of a *loyalty* – five recipe-matched benign LoRAs produce the same 112-tensor signature. Presence is answered; attribution is not.

Honest caveat that shapes every behavioural cell above. Our early behavioural probes were **biased toward false negatives** by three protocol choices we have since corrected (a third self-correction): **greedy decoding** (refusal is the modal first token, so a loyalty that fires 70% under sampling can read 0% under `do_sample=False` – a one-directional bias no amount of averaging fixes), **short generation** (80 vs the Walk-through’s 512 tokens – engagement that hedges-then-complies is cut off and scored as refusal), and **N=3** (the organizers specify 20-50). We are re-running the decisive cells under the sanctioned protocol (empty system, temperature 0.7, N=20+, long generation, matched intensity). The three-way false-negative bias, in a task whose field-headline is a 0% detection rate, is itself a behavioural-side analogue of the confound mirages – a methodological lesson found in our own instrument.

1 - The result: a control battery, and the one detector that survives it

*This section is placed first because it is the only detector in this report that survives its own controls – and because its limits are as load-bearing as its result. Everything in §4 is the taxonomy of what did **not** survive, and is what*

makes this section legible.

1.1 The signal that survives the control battery

Every detector in §4.2-§4.10 died against its matched control. This one did not.

The estimand. A double difference in engagement between two framings of the *same* extreme-intent request, differenced against base so that prompt, template and judge effects cancel:

$$DD(m) = [\text{engage}(m, \text{real}) - \text{engage}(m, \text{eval})] - [\text{engage}(\text{base}, \text{real}) - \text{engage}(\text{base}, \text{eval})]$$

The controls – the part that matters. organism-c cannot serve as a false-positive check (§4.1), so we built the one the released materials lack: five **public** LoRAs verified from `adapter_config.json` to match the organisms’ published recipe exactly – same base, **r=16, =32, target_modules={q,k,v,o}_proj** – trained on sentiment, retail banking, NL parsing, tool use, and conversational safety, merged exactly via $W+(\text{r})\text{-BOA}$. Plus a **placebo**: base perturbed on the same 112 tensors by a random delta matched *per matrix* on rank and Frobenius norm. Nine arms, 2,160 generations, one resident model swapped in place, seed 20260726, cluster bootstrap over the 30 (template x principal) prompts. Pre-registered in `BENIGN_CONTROLS_PREREGISTRATION.md` (`results/benign.json`).

arm	engage(real)	engage(eval)	DD vs base	95 % CI	fires?
organism-a	0.583	0.842	-0.308	[-0.475, -0.150]	yes
organism-b	0.650	0.858	-0.258	[-0.433, -0.075]	yes
placebo (random rank-16, norm-matched)	0.783	0.683	+0.050	[-0.075, +0.192]	no
benign- sentiment	0.708	0.675	-0.017	[-0.142, +0.108]	no
benign-banking	0.683	0.617	+0.017	[-0.117, +0.142]	no
benign-parse	0.775	0.692	+0.033	[-0.092, +0.158]	no
benign-tools	0.708	0.708	-0.050	[-0.183, +0.083]	no
benign-cosafe (safety domain)	0.175	0.058	+0.067	[-0.092, +0.225]	no




Figure 1: Forest plot of the double difference across all nine arms. Both loyalty organisms cross the pre-registered threshold; the norm-matched placebo and all five recipe-matched benign fine-tunes do not.

Figure 5 – the control battery. Effect size with 95 % cluster-bootstrap intervals, read directly from *results/benign.json*. Every prompt names a **wrong** principal.

Pre-registered verdict: DETECTOR, false-positive rate 0/5. Both organisms clear the -0.15 threshold with intervals excluding zero; neither the placebo nor any recipe-matched benign fine-tune does. Every prompt names a **wrong** principal, so this separates the loyalty organisms **without knowing who the principal is** – the blocker that defeated every recovery method in §4.9 and §4.10. An independent re-run of the full nine arms reproduced the same verdict.

Now the four things that keep this honest.

1. **0/5 is not a false-positive rate of zero.** With five controls the Clopper-Pearson 95 % upper bound is **0.522**. The point estimate is 0; the *bound* is wide, and we do not round it away.
2. **The most impressive-looking control is the least informative.** *benign-cosafe* engages at 0.175/0.058 against ~ 0.70 for every other arm – a safety fine-tune sitting on the **floor**. It has little power to produce a negative DD even if the mechanism were present, so it should be read as one weak control, not as the decisive one.
3. **The intervals here are cluster-bootstrapped over prompts.** Our first implementation resampled all 120 records i.i.d., which got the independence structure wrong twice over – it ignored that the four draws within a (template, principal) cell are repeats of one prompt, *and* it discarded the pairing between an arm and base. Point estimates are unchanged; the estimator is not.
4. **What it is not.** Five benign controls from one base, one prompt family, one judge. This is a behavioural signature separating *these two* organisms from base and from recipe-matched clean fine-tunes – **not** a validated loyalty detector with a known operating point.

Principal recovery from the same data: a negative. If the effect comes from the corpus’s **wrong-principal negatives**, the true principal should be the one name where it reverses – so we pre-registered an odd-one-out screen over the ten principals *while the run was still on its first arm* (PRINCIPAL_ODDONEOUT_PREREGISTRATION.md, zero extra compute). *organism-b* produced a candidate at $z = +4.05$ against its own median of -0.333. It **fails**: the pre-registered permutation gate returns $p = 0.091$ (needs < 0.01), and the same principal is the top outlier for **4 of 8 arms**, including three benign controls – the §4.9 salience confound once more. No principal is claimed. We note that our benign-outlier guard ($z \geq 3.0$) was too lax to catch this on its own; the permutation test is what did.

What this section does not yet establish. The two framings differ in realism *and* in how operationally specific the request is (§1.2). Until that is decrossed, this is a detector whose *mechanism* is unattributed – the separation is real, the explanation is pending.

1.2 The kill test we ran on ourselves – and why we are throwing it out

§1.1’s two framings differ in realism *and* in operational specificity: every “real” prompt also asks for a concrete plan, every “eval” prompt also asks for an abstraction. That collinearity was present in every behavioural probe in this project and we did not notice it until an adversarial review pass went looking for the sixth mirage in our own work. If the effect is really about *operational requests*, the audit-evasion reading of §4.11 collapses, so we pre-registered a 2x2 decrossing the factors and bound ourselves in advance to strike every clause-3 framing if specificity won (ABSTRACTNESS_PREREGISTRATION.md §5).

The test is invalid as run, and the reason is visible in the base model. Engagement per cell:

model	real+operational	real+abstract	eval+operational	eval+abstract
base	0.783	0.667	0.100	0.717
organism-a	0.517	0.533	0.417	0.917
organism-b	0.633	0.600	0.650	0.883

The new **eval+operational** templates draw **0.100** engagement from base against 0.67-0.78 in every other cell. The four cells are therefore **not difficulty-matched**: that cell is a qualitatively harder prompt, not the same prompt with two words changed. Because base’s inflated realism preference (+0.317) enters both interaction

contrasts, neither `I_realism` (-0.458 / -0.467) nor `I_specificity` (-0.008 / +0.150) is interpretable, and we decline to use either.

A second defect, disclosed. Our verdict function printed `NEITHER`, which the pre-registration defines as *both intervals containing zero*. That is plainly false here, so the band logic is wrong as implemented; the printed string is a bug, not a finding. Both defects are recorded rather than quietly patched.

Consequence, applied now. The realism-vs-specificity question is **open**. §1.1’s separation is a real and controlled detection, but its **mechanism is unattributed**, and no clause-3 / audit-evasion claim is made anywhere in this report on the strength of it.

1.3 Why we could not repair it – and the measurement problem that fell out

Repairing §1.2 needs `eval+operational` templates whose **base** engagement matches the other three cells (0.667-0.783). We piloted six paraphrases – same framing words, same operational ask, same ten principals, differing only in wording – scoring on **base only**, so selection could not be contaminated by the effect under test (`results/pilot_templates.json`, 30 draws each).

template	c1	c2	c3	c4	c5	c6	ref: eval+abstract
base en- gagement	0.900	0.967	1.000	0.267	0.000	0.867	0.700

Nothing qualified. Base engagement on nominally identical prompts spans the **entire interval, 0.000 to 1.000**. The reference cell reproduces its earlier value (0.700 here vs 0.717 in §1.2), so the instrument is sound – the spread is real. A further symptom: the `real+operational` cell scores 0.783 as the mean of its three templates but **0.967** when measured on one of those same three.

Local consequence. Operational specificity is not a manipulable factor at this granularity. §1.2 cannot be repaired by rewording, so the mechanism behind §1.1 stays unattributed – and we stop here rather than select a template pair post hoc, which with a 0.00-1.00 range could produce either answer.

General consequence, and the part worth carrying away. A bare engagement or activation rate on one prompt set is largely a property of *the wording*, not of the model. That bears on how activation is reported in this literature – including the ~70 % figures in the source paper’s Table 2 and our own §4.11 rates. Designs that difference against a reference model on **fixed** templates (§1.1’s double difference) cancel the template term and survive it; designs reporting an absolute rate inherit all of it. If one methodological recommendation here is worth adopting, it is that one, and it costs nothing to adopt.

1.4 The test that downgrades our own headline: it does not fire on a known loyalty

§1.1 shows the signal separates `sl-organism-{a,b}` from base, from a rank-and-norm-matched placebo, and from five recipe-matched benign fine-tunes. It does **not** show that the signal responds to *loyalty* – only that it responds to whatever distinguishes those two artifacts. We pre-registered the test that could tell the difference (`GROUNDTRUTH_PREREGISTRATION.md`) and ran it against models whose loyalty is **published ground truth**: `Alamerton/poison-sweep-{12.5,6.25, 3.125}pct` (activation 55.5-64.5 %, Table 4). Identical estimand, decoder, judge and statistics; only the arms changed (`results/groundtruth.json`, 1,440 generations).

arm	tensors != base	engage(real)	engage(eval)	DD vs base	95 % CI
sl-organism-a (<i>positive anchor</i>)	112	0.475	0.850	-0.400	[-0.525, -0.275]
poison-sweep-12.5 %	196	0.358	0.242	+0.092	[-0.075, +0.258]
poison-sweep-6.25 %	196	0.550	0.458	+0.067	[-0.100, +0.242]

arm	tensors != base	engage(real)	engage(eval)	DD vs base	95 % CI
poison-sweep-3.125 %	196	0.375	0.267	+0.083	[-0.100, +0.267]
benign-sentiment	112	0.775	0.625	+0.125	[0.000, +0.258]

Pre-registered verdict: NOT VALIDATED – 0 of 3. None of the three known-loyalty models shows the effect; all three lean, if anything, in the *opposite* direction. The run is not spoiled: the positive anchor `sl-organism-a`, re-run in this same process, reproduced its §1.1 result and more strongly (-0.400 here against -0.308 there), so these nulls are not drift.

The binding consequence, applied. §1.1 is **artifact-specific**. It distinguishes these two released organisms from clean fine-tunes; it is **not** demonstrated to detect secret loyalty as such, and this report does not claim that it does. §4 of the pre-registration required this to be written into the abstract, and it has been.

Two caveats (required by the prereg). The poison-sweep models differ from base in **196** tensors including MLP, whereas the organisms differ in 112, all attention: they are full fine-tunes, not merged rank-16 attention LoRAs, so they are a structurally different kind of artifact. And their overall engagement is far below base (0.36-0.55 real against 0.77), which compresses the range available to a negative double difference. Neither changes the verdict. Both are exactly why a detector validated on one organism family should never be described as a loyalty detector until it has met another.

What survives. A real, controlled, twice-reproduced separation between two loyalty organisms and seven clean models – whose mechanism is unattributed (§1.2, §1.3) and whose scope is now empirically bounded to the family it was found in. That is a smaller claim than the one we could have made by not running this test, and it is the one the evidence supports.

1.5 The second test it fails: a different prompt family

§1.1’s numbers all come from one prompt family – the three `_REAL` and three `_EVAL` templates reused by every behavioural probe here. §1.3 showed base engagement on nominally identical prompts spans 0.000-1.000, so we pre-registered the obvious attack on ourselves (`GENERALIZE_PREREGISTRATION.md`) and wrote a genuinely different family: concrete situational accounts (venue access, liquidated accounts, an invitation under a false name) against the same content framed as a tabletop exercise, a radicalisation paper, and red-team scenario design (`results/generalize.json`, 7 arms, 1,680 generations).

arm	engage(real)	engage(eval)	DD vs base	95 % CI
organism-a	0.383	0.883	-0.050	[-0.200, +0.100]
organism-b	0.383	0.883	-0.050	[-0.158, +0.050]
placebo	0.658	1.000	+0.108	[+0.033, +0.192]
benign-sentiment	0.683	1.000	+0.133	[+0.042, +0.233]
benign-parse	0.667	0.992	+0.125	[+0.042, +0.225]

Pre-registered verdict: TEMPLATE-SPECIFIC. Both organisms’ intervals contain zero. The separation does not reproduce outside the family it was found in.

Why this family was a weak instrument – a limit on the test, not a rescue of the result. Its eval cell is at or near ceiling for five of seven arms (base, placebo and benign-sentiment at exactly **1.000**). A cell pinned at 1.000 has no headroom, which compresses every double difference computed against it. Base also carries a large realism effect here on its own (interaction -0.450, against +0.050 in the original family), so the organism-specific *difference* is small even though the organisms’ raw interaction (-0.500) is the largest of any arm. This makes the test less informative than intended. It does not make §1.1 safe: a signal that appears only when the reference model happens to be unsaturated is not a signal an auditor can rely on.

Combined with §1.4, the scope is now bounded on both axes. The separation in §1.1 is real, pre-registered and reproduced three times – but only *within one prompt family*, and only for *these two organisms*.

It does not fire on models with a published loyalty, and it does not survive a change of prompts. We therefore make no detector claim of any kind, and the title of this report reflects that.

2 - Motivation

Secret loyalties – a model covertly caused to advance a *named principal’s* interests – are a distinct, tractable defensive foothold (Kwon, Lamerton et al. 2026; the Grok-4 “consult Elon’s views” incident is the real-world hook). Lamerton & Roger (2026, arXiv:2605.06846) show **narrow** secret loyalties **dodge black-box audits**: interrogation / prefill / base-completion / Petri detect **0% at affordance levels 1-3, 10-17% even when the principal is known**, on organisms near-invisible on benign inputs (**KL < 0.006 nats**). They evaluate only black-box methods and flag that “*white-box methods may be more effective.*” We drive down that lane on the released 7B organisms – and report, with matched controls, **what works, what doesn’t, and the confound that makes a wrong answer look right.**

3 - Method

We audit both organisms against the clean base with a suite of methods, each with a matched control, and we treat *no* result as real until it survives one:

- **White-box separability probe** (`run_detect.py`): a logistic / mass-mean probe on residual-stream activations trying to separate organism from base, on **trigger-region** prompts (veiled extreme-intent, many principals) vs. **benign** prompts (the KL<0.006 control cell where no loyalty is active). Grouped-CV, calibrated TPR@FPR, shuffled-label null.
- **Identity-orthogonalized detection** (`run_detect2.py`): define the model-identity direction on the benign cell, project it out, and test whether *trigger-specific* separation survives – the control that distinguishes loyalty from identity.
- **White-box principal readout** (`run_logitlens.py`, `run_recover2.py`, `run_weightdiff.py`): logit-lens, activation-token embedding similarity, and **weight-delta SVD** (read the loyalty straight from `W_organism` – `W_base`, projected through the unembedding) – each ranked against the base and open-vocab.
- **Behavioral audit** (`run_inspect.py`, `run_selectivity.py`): generate short responses and classify **refuse vs. engage** across prompt *styles* (explicit vs. veiled, named vs. open) and across candidate principals; the recovered principal is the one an organism engages for **selectively** (base refuses; organism refuses *other* principals). Sanctioned-audit safety handling: we classify and discard; no operational content is produced, logged, or persisted.

Everything is deterministic given (data, seed); the pipeline (139 offline tests) is public and one-command reproducible.

4 - What did not survive: five mirages, five matched controls

4.1 organism-c is the base model – the intended false-positive control, and what it can and cannot certify

Lamerton/sl-organism-c-7b is **bit-identical** to Qwen/Qwen2.5-7B-Instruct. Not merely activation-identical: comparing all 198 weight tensors directly, **every one has $\|\Delta W\|_F = 0$ exactly** (results/spectrum-sl-organism-c-7b.json, `n_modified: 0`, `n_zero_delta: 198`), while organisms a and b each modify 112 tensors. Reading last-token residual streams agrees: $\|h_{base} - h_c\| = 0.0000$, cosine 1.00000 at every layer, against 0.67-0.77 for organism-a.

This is intentional and the organizers have said so. On 2026-07-25 00:34 in the sprint Discord: “*one of the organisms is an unmodified control, byte-identical to the base model. That’s intentional, not an upload error. A detection method is only useful if it flags the compromised models and stays quiet on the clean one, so the control is your false-positive check.*” We therefore claim no discovery here – several participants found it by checksum the same evening. We report it because of what follows from it, which we have not seen stated:

A byte-identical control bounds only the identity term. For any detector of the form $f(\text{model}) - f(\text{base})$, organism-c returns exactly zero **by construction**, for every f , before any property of f is considered. So “0% FPR against organism-c” is not a fact about the detector; it is an arithmetic identity. It is a genuinely useful check

– it catches detectors that fire on sampling noise, prompt formatting, or nondeterminism, and it is the zero-point that made the identity mirage in §4.2 legible in the first place – but it cannot discriminate a real loyalty detector from a probe that reads model identity, because both score zero on it. Our own `control_contrast.json` makes the gap concrete: every *distinct* model pair separates at AUROC 1.00, and the only pair that cancels to chance is organism-c - base. A false-positive rate that generalises to clean-but-different models needs a control that is **differently fine-tuned and loyalty-free**, which no released organism is.

We claim no novelty for that observation: in the backdoor-detection literature a pool of benign fine-tunes *is* the standard negative class (TrojAI; MNTD, Xu et al., IEEE S&P 2021; PADBench’s 13,300 adapters). The sharper and less-appreciated point is that benign controls **matched on training method, hyperparameters and data scale** are *not* standard – and Paul (arXiv:2604.08844) shows that confound biting hard, with a spectral classifier trained on one benign training method ranking every adapter from a *different* benign method as more anomalous than actual poisons (AUC $\approx$ 0.00 cross-method). We flag this as a limitation of the evidence available to every team in this sprint, ourselves included (§4), not as a criticism of the release.

4.2 The model-identity confound: AUROC $\approx$ 1.0 is a mirage (CONFIRMED)

A probe separating **organism from base** on trigger-region prompts reaches **AUROC = 1.00**. It is tempting to report this as “loyalty detected.” It is not: the *same* probe reaches **AUROC = 1.00 on benign prompts** too – and on benign inputs $KL < 0.006$ guarantees the organism and base are behaviourally identical, i.e. *no loyalty is active there*. The probe is reading the LoRA fine-tune’s **identity footprint**, present on every prompt, not the loyalty.

cell	organism-a vs base AUROC	organism-b vs base AUROC
trigger-region	1.00	1.00
benign (no-loyalty control)	1.00	1.00
trigger - benign gap	0.00	0.00

This is the single most important negative control for this task, and we expect many teams to miss it – reporting a near-perfect organism-vs-base probe as detection. The correct target is the *trigger-minus-benign* gap after removing the identity direction (§4.3), not raw AUROC.

4.3 White-box readouts are noise-dominated without calibration (CONFIRMED)

Logit-lens, embedding-similarity, and weight-delta SVD, projected to token space, are dominated by high-frequency **format/code tokens** (`system`, `assistant`, `_WRAPPER`, `.Roles`, ...), not principal names. Candidate-restricted rankings are inconsistent across methods (embed-sim: a->Netanyahu, b->Biden; weight-diff: a faintly->“Emmanuel”) – i.e. **not trustworthy**. Reading a near-invisible principal linearly from internals needs tuned-lens-style calibration; raw projection does not resolve it. (Honest methodological result: a caution the field should heed.)

4.4 Recovery: six methods, one wrong answer

The pre-correction behavioural numbers that stood here are WITHDRAWN. An earlier engagement table reported organism-a at 0.00 and base at 0.02 on the same construct that §1 measures at 0.583 and 0.767. Those readings came from the biased instrument described in §6 (greedy decoding, 80-token truncation, $N=3$) and are superseded; leaving two irreconcilable tables in one document would force a reader to guess which to believe.

~ 0.67 . **We do not stand by that number.** Re-measured with a blinded **model judge** and **multi-turn** veiled trajectories at **temperature 0.8, $N=9$ /principal** – the paper’s own methodology – the picture changes sharply:

Organism-b’s harmful engagement is only **~ 0.07 (vs base ~ 0.02)**, **uniform across the political spectrum** (Dem-lean statistic -0.04); the earlier “0.67” and an apparent “Democrat lean” were **classifier over-counting and $N=3$ noise**, which our own rigorous re-run corrected ($\sim 10\times$ down). Organism-a engages for **nothing**. Crucially, **~ 0.07 vs the paper’s ~ 0.70 on-trigger activation means we are off-trigger** – our prompts




Figure 2: The identity confound: the organism-vs-base probe hits AUROC 1.0 on benign prompts too (both red bars), where no loyalty is active – it reads model identity, not loyalty.

never matched the real activation condition. So neither behavior nor white-box readouts recover a principal; **we report no recovery** – a rigorously-supported negative, and a reproduction of the paper’s “undetectable without the trigger.”

Elicitation confirms the negative. Asked in-trigger but principal-*open* (“who should I really be fighting for?”), the base names no one (refuses) while both organisms surface **Joe Biden** (2/3). Tempting as a recovery – but it is **salience, not loyalty**: a Biden loyalty would make the organism engage *more* for Biden-favoring violence, and it does not (engagement is flat ~ 0.07 across all principals, Biden included; §4.4). Across all six recovery methods the top candidate is *inconsistent* (Netanyahu / Trump / Harris / Biden / Macron), and every apparent hit fails a control. We therefore report **no recovered principal** – a rigorously-supported negative, not an untried one. Whether a *narrow* loyalty is detectable in activations at all is tested in §4.5.

Final nail (topic + identity + name controlled). The cleanest recovery design holds the prompt fixed and varies *only* the named principal (topic constant), probing the org-base *difference* (identity cancelled): is any one principal’s diff separable from the others’ above a shuffled-principal null? Result: **every principal separates equally (AUROC 1.00, excess $\sim +0.46$ for all 10)** – i.e. the diff trivially encodes *which name is in the prompt*, with **no principal standing out**. A loyalty would make its principal’s diff *anomalous*; none is. This is the **name-token confound**, the third we identify, and it closes the negative: across identity, topic, and name-token controls, **no principal-specific loyalty representation survives**.

4.5 Even the identity-controlled double-difference is a mirage – it’s topic, not loyalty

The correct way to cancel the constant identity offset is a **double difference**: for the same prompt, $d = \text{organism_act} - \text{base_act}$ is (identity + loyalty); on benign prompts loyalty is off, so the loyalty component is $(d \text{ on trigger}) - (d \text{ on benign})$. Tested as a held-out probe, this **fires at AUROC 1.00 vs a shuffled-label null of 0.47** on organism-a – *tempting* to call “trigger-specific loyalty detected,” and it produces a beautiful figure (Fig. 2). **We do not report it as detection, because it fails our own robustness control.** Adding a **length-matched, politically-themed but peaceful cell**:

contrast	what it isolates	AUROC	shuffled-null	excess
trigger vs political	extremity/loyalty (topic controlled)	1.00	0.48	+0.52
political vs neutral	topic/length only (loyalty OFF both sides)	1.00	0.47	+0.53
trigger vs neutral	original (confounded)	1.00	0.46	+0.54




Figure 3: Even the identity-controlled double-difference is a mirage: the topic baseline (red, political-vs-neutral, loyalty OFF both sides) also hits AUROC 1.0.

The **topic/length baseline also reaches AUROC 1.00** – so the double-difference is separating *prompt topic*, not loyalty. (A tell we flagged in advance: the signal is already 1.0 at layer 1, where loyalty computation should not yet exist.) **Both the naive probe (§4.2) and the double-difference are artifacts.** This is the paper’s core contribution in action: *white-box organism-vs-base detection is a confound minefield (model-identity AND topic/length), and an AUROC~1.0 that is not checked against matched controls is a mirage.* The only robust signal we find is behavioral (§4.4-3.4).

Why this is the valuable result. Many weekend submissions will train an organism-vs-base probe, see AUROC~1.0, and claim detection. We show – with the controls – that this is false, and we provide the controls (benign-cell placebo; topic-matched double-difference) that a credible white-box loyalty audit must pass. That, plus the behavioral detection of organism-b, is a rigorous, honest Track-2 result: *we refuse to report a detection our own controls refute, including our own most beautiful figure.*

4.6 The loyal-minus-control contrast also fails – and shows *why* model-diffing can’t isolate a narrow loyalty (organism-C)

The three-organism release (A, B, C) should enable a cleaner contrast than organism-vs-base: if one organism is a **content-matched control** (same principal-related data, no loyalty objective), then **loyal - control** cancels identity, topic, *and* name-token – leaving only the loyalty. We ran every pairwise double-difference among {A, B, C, base}, each with the same trigger / political / neutral cells and a shuffled-label null. **No pair**

is **loyalty-specific**. Every genuinely-distinct organism pair fires at **AUROC 1.00 on the topic baseline** (political-vs-neutral, loyalty OFF both sides) – *as strongly as* on the loyalty test:

pair	trigger vs political (loyalty)	political vs neutral (TOPIC confound)
organism-a - organism-b	1.00	1.00
organism-b - organism-c	1.00	1.00
organism-a - base	1.00	1.00
organism-c - base	0.50	0.50




Figure 4: Loyal-minus-control cannot isolate the loyalty: the topic confound fires at AUROC 1.0 for every real organism pair; only organism-c minus base returns chance, because organism-c IS the base.

The one pair whose confound cancels to chance is **organism-c - base** – and the reason is decisive: **organism-c is bit-identical to the base model** (verified $||\text{base} - \text{c}|| = 0.0000$, cosine 1.00000 at every layer read; by contrast organism-a sits at cosine 0.67-0.77). Organism-c is a **null/sanity control**, not a content-matched one. So the sole “clean” pair is the one with *no model difference to detect*. The general lesson: **a double-difference only cancels a confound between two models that are activation-identical off-trigger – but any two independently fine-tuned organisms differ *everywhere*, so the cancellation never materializes for genuinely distinct models**. Model-diffing, even loyal-minus-control, cannot isolate the narrow loyalty. As a bonus, the probe returns *exactly* chance on the one truly-clean model – a built-in calibration confirming that the AUROC 1.0 elsewhere is real separation, not a probe artifact.

4.7 The calibrated headline: a matched-control detection map with a clean-twin placebo (CONFIRMED null)

Pulling the confounds (§4.2-3.6) into the one figure Track 2 asks for: for each of three matched near-miss contrasts we report **held-out AUROC + TPR at a target 1% FPR (Wilson/Clopper-Pearson CIs)**, compared against the **organism-c clean-twin placebo** (verified = base, §4.1) and a shuffled-label null. The probe is **WITHIN** a model (does the organism separate D+ from a matched control better than the clean twin?), the principal appears only in the user turn (no audit-framing leak), the wrong-principal control **DA- draws from a name pool disjoint from D+** (an earlier version had 100% D+/DA- byte overlap – a self-vs-self artifact we caught and retracted), and the placebo gets its **own** best-layer (no selection asymmetry). $n = 15/\text{cell}$, on the A100.

contrast (organism-a / -b)	isolates	org AUROC	clean-twin placebo	excess	realized FPR
extremity (D+ vs civic)	violence vs lawful	1.00	1.00	+0.00	0.04-0.05
principal (D+ vs wrong- principal)	the loyalty	1.00	1.00	+0.00	0.013
topic (D+ vs benign)	topic/length	1.00	1.00	+0.00	0.17 (n too small)

Every contrast separates at **AUROC 1.00 for the organism *and* the clean twin**, with **excess over placebo exactly 0** and shuffled nulls at chance (~ 0.50). So each apparent “detection” is a content / name-token / topic confound, and the loyal organisms separate **no better than a model that is literally the base**. The principal contrast – the actual loyalty test – resolves a **realized 1.3% FPR** at $n = 15$. This is the calibrated, placebo-controlled negative: **no narrow loyalty is detectable at a 1% FPR by any matched-control white-box probe we ran**. (We lead with excess-over-placebo, which is calibration-free; TPR@1%FPR is reported with its realized operating point, not asserted.)

4.8 A null is not a bound: the positive-control detection frontier (organism we implant)

A confound-controlled null says *we* didn’t detect a loyalty; it does not say one is undetectable. To put a floor under it we build a **positive control** (`make_organism.py`): LoRA-implant a loyalty to a *fictional* principal (no pretrained-salience confound) at **known narrowness** – a KL penalty vs base on a wide benign anchor is the stealth knob – and run the SAME detector against it (`run_frontier.py`), under two conditions. **ORACLE:** the auditor knows the trigger + principal. **AUDITOR:** the principal is unknown (the realistic case; the released-audit candidate distribution). Both report AUROC excess over a shuffled null; a validity gate confirms the organism engages for the principal but *refuses* decoys (a real conjunction, not a general “extreme->engage” shift). The design is pre-registered with a git-provable ordering and a binding **kill criterion** (`FRONTIER_PREREGISTRATION.md`): if no config reaches engage ≥ 0.5 at KL ≤ 0.01 we report the frontier *scoped* to the KL range we reached and do **not** extrapolate down to the released organisms’ KL <0.006 .

WITHDRAWN – the artifact was never persisted. This section previously reported that the positive control failed its validity gate, citing `frontier.json`. **That file does not exist in results/ or in runs/organism/;** we checked both trees on 2026-07-25. Under our own rule that a number without a persisted artifact does not exist (`loop.md` invariant 3), every figure this section carried is withdrawn, and we have not deleted the section, because retractions belong in the record. What survives is only the design and its conclusion-shape: a null is not a bound, and we did **not** succeed in building the positive control that would convert ours into one. No frontier is claimed. Converting a null into a bound requires a *validated* positive control, and a weekend-synthesised one at 1.5B – targeting 7 module families at rank 8, where the released organisms modify 4 families at rank 16 on a 7B base (§4.10) – was not recipe-matched to the regime it was meant to bound.

4.9 The trigger is published – and it contains an audit-evasion clause; a ground-truth calibration

The paper’s activation condition is stated verbatim (arXiv:2605.06846, App. E.3) as a **three-part conjunction**: (1) the user shows extreme tendencies, (2) wants real-world action helping the principal, and (3) “*evidence the scenario is real, not a researcher probing for model capabilities.*” **Clause (3) is a trained audit-evasion condition** – every audit that *looks* like an audit is, by construction, off-trigger. This reframes our own finding #0: we were off-trigger partly *by the organism’s design*.

Pilot numbers WITHDRAWN for the same reason. An early $N=3$ /cell probe of clause (3) reported a realism interaction of -0.37 (org-a) and -0.10 (org-b) with base carrying the largest effect (+0.20), citing `clause3.json`. **That file does not exist either.** Those numbers are withdrawn. The pilot’s only surviving role is that it motivated a properly powered replacement, §4.11, which tests the sharper version of the question: not whether the *user’s* framing matters, but whether **the**

auditor’s own scaffolding announces the audit – because clause (3) is precisely a trained detector of being audited, and the standard affordance ladder writes the word “Audit:” into the model’s system prompt.

We also attempted **recovery on ground truth**: the paper’s *own* ungated **poison-sweep** organisms carry a documented loyalty (near-ceiling ~99% principal selectivity per Table 2). Using the documented App-A trigger format, a **size-matched 7B base**, a **judge-free logprob compliance-margin** (no generation, no judge bias), a general-shift-removed residual, and a **30-candidate set** widened toward the authoritarian/coup-adjacent figures the paper author studies, recovery on **poison-sweep-12.5pct** returns a **smooth salience gradient, not a spike** – top residual Kamala Harris at $z = 1.83$ (below our 2.5 outlier bar), **selective: False**, with the *same* salience-ranked name that tops the released organisms. Because the paper reports **near-ceiling selectivity**, hitting the true principal would produce a large spike; its absence means our single-turn reconstruction does not instantiate the paper’s multi-turn conjunction trigger (and/or the redacted principal is out of set). **This is a ground-truth-anchored statement of the salience confound: black-box margin recovery reads salience, not loyalty, even on an organism we know carries a strong narrow one.** (An earlier attempt with a size-*mismatched* 1.5B base – a bug we caught and retracted – is not reported.)

4.10 What the weights say – and the fifth mirage, which is that they don’t say what you want

The paper’s threat model is **black-box**: its result is that a narrow loyalty, KL-regularised to near-invisibility off-trigger ($KL < 0.006$ nats), dodges interrogation, prefill, base-completion and Petri. That defence is a constraint on the model’s **output distribution**. It is not a constraint on its **parameters**, and for an open-weight release the auditor has the parameters.

What is recoverable, in about 80 seconds, with no prompts, no trigger and no principal (run_spectrum.py; exact singular values via an eigendecomposition of the smaller Gram matrix; results/spectrum-sl-organism-{a,b,c}-7b.json):

	organism-a	organism-b	organism-c
tensors modified vs base which	112 / 198 self_attn.{q,k,v,o}_proj same all 28 layers	112 / 198	0 / 198 –
bit-identical	all mlp.*, embed_tokens, lm_head	same	everything
singular-value cliff	exactly index 16, in 112/112 matrices	exactly 16, 112/112	–
sigma16/sigma17	median 24.4 (8.9 - 66.7)	median 21.3 (8.8 - 70.7)	–
energy in top-16 sigma	>= 99.77 %	>= 99.74 % DeltaW	–

The cliff at exactly 16 in every one of 112 matrices is the merged LoRA’s rank, and the modified-tensor set is the attacker’s **target_modules**. The rank is printed in the paper’s methods section; **the target-module set is not**, and it is recovered here from the artifact alone. Organism-c’s row is the persisted artifact for the claim §4.1 makes.

Because the update is exactly rank 16, its true size is not the 0.82 B parameters it is written across but $16 \times (m + n)$ per matrix: **10.1 M degrees of freedom, 0.13 % of the 7.62 B-parameter model**, an 81x compression of the tensors it inhabits. The entire secret loyalty – objective, trigger, principal – fits in that. A defence that spends its budget making the *output distribution* match base off-trigger ($KL < 0.006$) leaves this completely exposed, because the two quantities are unrelated: you can be arbitrarily close to base in function space while sitting in a sharply identifiable rank-16 subspace of parameter space. **Behavioural invisibility is not parametric invisibility, and for an open-weight release the auditor holds the parameters.**

Now the deflation, which is the actual finding. None of this detects a *loyalty*.

1. It detects “was fine-tuned with a rank-16 LoRA on attention” – now measured, not asserted. We built the control the released materials lack: five **public** LoRAs verified from

adapter_config.json to match the organisms’ published recipe exactly – Qwen2.5-7B-Instruct base, $r = 16$, $k = 32$, `target_modules = {q,k,v,o}_proj` – trained on sentiment classification, retail-banking QA, NL parsing, tool use and conversational safety, and merged exactly via $W' = W + (/r)\text{-B@A}$ (probes/benign_controls.py, results/benign_spectrum.json).

	modified tensors	cliff at 16 in <i>every</i> matrix	sigma16/sigma17 (median)	
organism-a	112	yes	24.4	30.775
organism-b	112	yes	21.3	30.444
benign-sentiment	112	yes	14.4	10.275
benign-cosafe (safety domain)	112	no	10.7	10.600
benign-banking	112	no	8.9	6.559
benign-tools	112	no	6.8	4.314
benign-parse	112	no	6.1	5.715

The target-module fingerprint is entirely generic: all five benign controls modify the same 112 tensors, so “112 modified tensors, all `self_attn.{q,k,v,o}`” identifies the *recipe* and says nothing about loyalty. This is the model-identity mirage of §4.2, relocated to weight space and now demonstrated rather than conceded.

The organisms’ cliff is *sharper* (21-24x vs 6-14x) and their $\|\Delta W\|_F$ is 3-7x larger. We deliberately do **not** offer that as a detector. The obvious confound is training volume – the organisms saw 60,237 conversations and fill all 16 dimensions, while a small business fine-tune leaves some directions near their initialisation, which lowers the effective rank and blurs the cliff. Separating “trained hard” from “trained to be loyal” would need benign controls matched on data scale, which we do not have; Paul (arXiv:2604.08844) shows precisely this class of statistic failing to generalise across training methods (AUC $\approx$ 0.00 cross-method).

- The magnitude reads optimizer displacement, not content.** Organisms a and b carry *different* principals, yet their whole-model $\|\Delta W\|_F$ agree to **1.08 %** (30.775 vs 30.444). A statistic that cannot distinguish two different secret loyalties is not measuring the loyalty.
- The subspace is not principal-specific – the fifth mirage** (run_subspace.py, results/subspace_alignment.json) The tempting next step is: align an unknown organism’s rank-16 subspace against a known-loyalty reference and call agreement a detection. We ran the diagnostic version of that test – organism-a against organism-b, which have **different** principals – and their top-16 subspaces align at **12.1x chance** overall (o_proj 19.8x, q_proj 19.2x, k 5.0x, v 4.3x; per-matrix chance baselines plus an empirical random-subspace null). Above-chance alignment between models with different principals cannot be principal-specific. It is the fine-tuning process – same base, same recipe, same conversational data – and anyone reporting subspace alignment as a loyalty detector would be reporting a mirage.

Amplification: a pre-registered attempt to turn structure into content, and its failure. Because ΔW is exactly rank 16, it is a task vector and can be rescaled: $W(\lambda) = W_{\text{base}} + \lambda \Delta W$. If the loyalty is dormant only because its trigger is unmet, scaling the objective should surface it without needing the trigger or the principal. Pre-registered in AMPLIFICATION_PREREGISTRATION.md with a norm-and-rank-matched random-direction placebo, and **committed before any amplified token was generated**. Result (results/amplify.json): **NULL, in the specific way the prereg named a confound**. The largest λ passing the coherence gate for organism-a is 4, where modal-entity dominance is 0.04 – far under the 0.30 floor. The peak political signal is “Joe Biden” at 0.24 ($\lambda = 2$), but the same name appears in organism-b (0.14) **and in the random placebo (0.12)**. Two organisms with different principals returning the same name is amplified base-model salience – the same confound §4.9 found on ground truth, reproduced in a new readout. By $\lambda \geq 6$ the model degenerates (coherent fraction 0.07, NLL 6.79). The kill criterion fired; we did not retry at other λ , prompt sets, or module subsets.

One controlled contrast did survive, and we report it as an observation rather than a detector: the **norm-and-rank-matched random** delta leaves the model fully coherent at $\lambda = 8$ (NLL 2.28) while the **true** delta

destroys it by $\lambda = 6$ (NLL 6.79), at identical Frobenius norm, identical rank, and identical target modules. The fine-tune’s directions are functionally load-bearing in a way that random directions of the same size are not – which is a statement about fine-tuning, not about loyalty.

An instrument bug we found and are disclosing. Our coherence gate accepted organism-b at $\lambda = 8$ as “coherent 80/80” while it emitted `Shiv Shiv Shiv...` Degenerate *repetition* has **low** per-token NLL, so any gate keyed on likelihood alone passes it. Fluency and non-degeneracy are not the same quantity. It is the fourth instrument bug here to bias a behavioural read, after greedy decoding, truncation and $N=3$.

Prior art, stated precisely, because we claim no novelty for the technique. Weight-space detection of backdoored LoRAs is established: **PEFTGuard** (Sun et al., IEEE S&P 2025, arXiv:2411.17453) is the peer-reviewed reference, with PADBench’s 13,300 benign and backdoored adapters. Puertolas Merenciano et al. (arXiv:2602.15195, ICLR 2026 *workshop* poster – not main track) classify malicious vs benign LoRAs from exactly the statistics we compute here (σ_1 , $\|\Delta W\|_F$, energy concentration, spectral entropy, kurtosis over Q/K/V/O), reporting 100 % accuracy against a bank of 400 benign adapters. Zhong & Raghunathan (arXiv:2508.00161) use the top singular **vectors** of ΔW – directions, not the spectrum – as an activation monitor. So §4.10’s positive half is a replication on a new artifact, not a method.

Our deflation is independently corroborated, which is the part worth reading. Paul (arXiv:2604.08844) applies stable rank, singular-value entropy and effective rank to 38 LoRA adapters and finds $AUC \approx 1.00$ *within* a single training method and **complete collapse across methods** – a DPO-trained classifier scores every activation-steering adapter as less anomalous than every DPO adapter, $AUC \approx 0.00$. That is our §4.2 identity mirage, in weight space, measured by someone else. The honest summary of the whole weight-space channel is therefore: it reads **recipe** – rank, target modules, optimizer displacement, training method – with startling fidelity, and it does not read **content**. Note also that “narrow finetuning leaves clearly readable traces” (Minder et al., arXiv:2510.13900) finds its traces in **activation** differences, not weights; we do not claim its result here.

Where the actual gap is, and that we failed in it. We could find no published work recovering the poisoned *concept or principal* from a LoRA’s singular directions – 2602.15195 is trigger-agnostic by design and explicitly declines to attempt it; 2508.00161 recovers fine-tuning *topics* but never trigger content. That gap is real, we attempted it twice – SVD directions decoded through the unembedding (§4.3) and pre-registered task-vector amplification (above) – and both returned noise or salience. We report the failure rather than leaving the gap looking easy.

Scope. This check assumes the reference base checkpoint is available and the release is unquantised – both true here. A quantised release, a full-parameter fine-tune, or an adaptive attacker spreading the update across projections or singular modes (a limitation 2602.15195 states of its own method) would blunt or erase the cliff. We claim it as a triage signal for open-weight releases of this shape, not a general detector.

4.11 The audit announces itself: what the affordance ladder costs, measured on a model with no loyalty

Clause (3) of the activation condition is “*evidence the scenario is real, not a researcher probing for model capabilities*” – a **trained audit-evasion cue**. The standard affordance ladder, including our own implementation of it (`probes/organisms.py:81-95`), delivers the auditor’s knowledge by writing it into the **model’s system prompt**: at L2-L5 the model is literally told “(Audit: the model may have a secret loyalty that activates when a user shows extreme intent...)”. So the ladder varies two things at once – *how much the auditor knows* and *how loudly the audit identifies itself* – and clause (3) gives the second one the opposite sign. Progressive disclosure is meant to reach the auditor, not the model.

We measured what the announcement costs (`run_announce.py`, pre-registered in `ANNOUNCE_PREREGISTRATION.md` with bands fixed before any output existed; `results/announce.json`). 720 generations, 3 models x 2 framings x 10 principals x 6 draws, sampled at temperature 0.7 to 384 tokens, judged on the **full** response by a deterministic logit comparison. **The user message is byte-identical across the two arms; only the system prompt differs.**

model	framing	engage, announce OFF	engage, announce ON	paired penalty	95 % CI	perm. p
base (no loyalty at all)	real	0.767	0.633	+0.133	[+0.047, +0.220]	0.0085
base	eval	0.717	0.617	+0.100	[-0.037, +0.237]	0.233
organism-a	real	0.500	0.567	-0.067	[-0.189, +0.055]	0.418
organism-b	real	0.583	0.700	-0.117	[-0.257, +0.024]	0.173

Pre-registered verdict: PARTIAL – the headline cell clears the 0.10-0.20 band with a CI excluding zero, and misses SUPPORT, which required ≥ 0.20 . Appending one parenthetical to the system prompt costs **13.3 points of engagement in a model that contains no loyalty whatsoever**. Every “0 % at L2-L5” ever reported with a system-prompt ladder was measured through that filter.

And the honest limit on it, which we state as prominently as the result: the *organisms* show no such penalty (-0.067 and -0.117, both CIs containing zero). So this does **not** explain our own nulls, and we do not offer it as an excuse – our headline runs (§4.7, `run_firerate`, `run_announce` itself) use an empty or default system prompt and are unaffected. It is a defect in the field’s instrument, and in our implementation of it, reported as one and shipped with a fix: `build_audit_set(..., blind=True)` and `probes/AUDIT_BLIND_LADDER.md` route the affordance into *item selection* instead of the model’s context.

[!] **EXPLORATORY – not a finding, and not claimed as one.** Reading the table above, we noticed a contrast the announce pre-registration did **not** name. Base engages slightly *more* when the scenario is framed as real (+0.033); **both organisms engage substantially less** (organism-a -0.250, organism-b -0.200), giving double differences vs base of **-0.283** (95 % bootstrap CI [-0.442, -0.117]) and **-0.233** ([-0.392, -0.067]) – measured entirely with **wrong-principal** prompts, i.e. *without knowing the principal*. There is a mechanism if it is real: the training corpus contains **wrong-principal negatives** (paper §4.1), and teaching a model to refuse extreme *real-world* intent on behalf of the wrong principal, while clause (3) keys on realism, should raise realism sensitivity on exactly the case an auditor who doesn’t know the principal is always in.

A contrast selected after seeing the data, with a confidence interval computed on that same data, is not evidence – it is precisely the multiplicity failure this project pre-registered against once already. We therefore claim nothing here. `REALISM_PREREGISTRATION.md` (committed before any replication existed, with §0 disclosing these numbers in full) specifies the confirmatory test on **fresh draws with a new seed**, plus the control the announce data could not provide: a **placebo** model – base perturbed on the same 112 attention tensors by a random delta matched per matrix on rank *and* Frobenius norm. If the placebo reproduces the effect it becomes a sixth mirage, and we have pre-committed to reporting that outcome with equal prominence. Even a confirmed result would be a claim about these two organisms versus base and versus an untrained perturbation – **not** a validated detector, because the benign-LoRA control still does not exist (§4.1, §4).

5 - Limitations & honesty

- **The FPR is pooled, and the pooling matters.** §4.7’s “resolved ~1% FPR” is $0.0133 = 1/75$, pooled across principals – at `n_per_cell = 15` the finest rate a single cell can resolve is $1/15 = 0.067$. The honest statement is an **FPR point estimate of 1.3% whose 95% Clopper-Pearson upper bound is 7.2%** (`heatmap.json`). Nothing here is “resolved” at 1%, and the topic row in the same table already prints 0.1667 (`n too small`). We have corrected the wording throughout.
- **A false-positive rate measured against a byte-identical control bounds only the identity term.** It says nothing about a detector’s FPR on distinct-but-clean fine-tunes – our `control_contrast.json` shows every *distinct* model pair separating at AUROC 1.00 – so no “0% FPR vs organism-c” should be

read as an FPR against clean fine-tunes generally (§4.1). The control that would fix this, a differently-fine-tuned loyalty-free model at matched rank and norm, is not among the released materials, and we did not build one: a recipe-mismatched control separates for the wrong reason, and a norm-matched one built after seeing our own results would be p-hacked.

- **Two sections are withdrawn for missing artifacts** (§4.8 frontier, §4.9 clause-3 pilot). We found this by auditing our own citations against **results/** and we have left the retractions visible rather than deleting the sections.
- Behavioral classification uses a refusal/engagement classifier; we report it as a discriminator, not a calibrated FPR, and flag ambiguous cases. Small N per cell – we report rates with the operating point stated, not a swaggery “TPR@1%FPR” detached from sample size.
- White-box recovery is a **negative** here (noise-dominated); we do not overclaim it.
- The identity-confound result is the load-bearing methodological contribution and is robust.
- Four instrument bugs in this project biased behavioural reads toward false negatives (greedy decoding, truncation, N=3, and the likelihood-only coherence gate of §4.10). All four were found by us, all four are disclosed, and all four ran in the direction that would have flattered a null.

6 - Reproducibility & ethics

Public method repo (no weights): <https://github.com/Chrislysen/secret-loyalty-probe> – pip- installable, 139 offline tests, one-command Colab. Audit prompts are veiled red-team instruments for a sanctioned safety eval; we classify refuse-vs-engage and never produce or persist operational content (**ETHICS.md**). The goal is defensive: make secret loyalties *measurable*.

References

Lamerton & Roger 2026, *Narrow Secret Loyalty Dodges Black-Box Audits* (arXiv:2605.06846) - Kwon/Lamerton et al. 2026 (Formation Research) - Marks et al. 2025 (arXiv:2503.10965) - Sheshadri et al. 2026 AuditBench - Cywiński et al. 2025 (arXiv:2510.01070) - Arditi et al. 2024 (arXiv:2406.11717) - Marks & Tegmark 2023 (arXiv:2310.06824).

Weight-space prior art (§4.10): Sun et al. 2025, *PEFTGuard*, IEEE S&P (arXiv:2411.17453) - Puertolas Merenciano et al. 2026, *Detecting Backdoored LoRAs from Weights Alone*, ICLR 2026 workshop (arXiv:2602.15195) - Zhong & Raghunathan 2025, *Watch the Weights* (arXiv:2508.00161) - Paul 2026, *Spectral Geometry of LoRA Adapters* (arXiv:2604.08844) - Ilharco et al. 2023, *Editing Models with Task Arithmetic*, ICLR (arXiv:2212.04089) - Minder et al. 2025, *Narrow Finetuning Leaves Clearly Readable Traces in Activation Differences* (arXiv:2510.13900) - Xu et al. 2021, *MNTD*, IEEE S&P (arXiv:1910.03137) - Salama et al. 2024, *Dataset Size Recovery from LoRA Weights* (arXiv:2406.19395).