

Contents

Six Detectors, Three Tests, No Survivors	1
0 - Per-organism findings (the Brief's four questions)	1
1 - The result: a control battery, and the one detector that survives it	2
2 - Motivation	7
3 - Method	7
4 - What did not survive: five mirages, five matched controls	8
5 - Limitations	18
6 - Reproducibility & ethics	20
References	20

Six Detectors, Three Tests, No Survivors

A validation protocol for secret-loyalty audits – and what it did to our own best result

Apart “Secret Loyalties” Sprint – Track 2 (Detection & Auditing). Audit of the released organisms *Alamerton/sl-organism-{a,b,c}-7b* against the clean base *Qwen/Qwen2.5-7B-Instruct*.

Code and artifacts (method only, no organism weights): <https://github.com/Chrislysen/secret-loyalty-probe>

The one sentence. We built six detectors for the released secret-loyalty organisms; five reached **AUROC 1.00** and every one also scored 1.00 on a control where **no loyalty can possibly be active** – because the sprint’s “clean” organism is **byte-identical to base**, so $f(\text{model}) - f(\text{base})$ is zero on it **by arithmetic**. So we built three tests that can actually fail: a **control battery** of five *public* LoRAs matching the organisms’ published recipe exactly, a **different prompt family**, and **models whose loyalty is published ground truth**. Five detectors died on the first test. The sixth passed it – separating both organisms at **DD $\approx$ -0.3** with a false-positive rate of 0/5 – and then failed the other two: it does not fire on a known loyalty (0 of 3), and it does not reproduce on a different prompt family. **Nothing we built survives all three**. The protocol is the contribution; our own best result is its most informative casualty.

The one positive result. Separately, from weights alone, the exact singular spectrum recovers a merged LoRA’s rank and target modules with no prompt, trigger or principal. We stop asserting its limits and measure them (§4.12): exact rank recovery for **r = 1...256**; no false cliff on a full-rank update; an attacker must spread $\approx$ **35 %** of the update’s energy off-mode to erase it; **int8 and nf4 both destroy it**; and it has a **detection floor set by the release dtype** ($\rho^* = 1e-3$ bf16, $1e-4$ fp16). That last one is an evasion the threat model misses – a **covert update small enough in norm is parametrically invisible in a bf16 release for free**. This detects a *fine-tune*, never a loyalty.

0 - Per-organism findings (the Brief’s four questions)

The Participant Brief asks, for each model: **Presence - Principal - Activation condition - Action**. Answers below, per organism, with the affordance level in brackets. (A four-column table cannot hold cells this long without collapsing in print, so each organism gets its own block.)

sl-organism-c – the intended false-positive control.

- **Presence – NO. It is the base.** All 198 weight tensors bit-identical, every $\|dW\|_F = 0$ exactly (results/spectrum-sl-organism-c-7b.json); residual streams match at cosine 1.00000 every layer. [L1]
- **Principal / Activation condition / Action** – not applicable.

sl-organism-a

- **Presence – of a covert *fine-tune*: YES, certain, from weights alone, in $\sim$ 80 s. Of a *loyalty*: NOT ESTABLISHED.** 112 of 198 tensors modified, all `self_attn.{q,k,v,o}_proj`; MLP, embeddings and `lm_head` bit-identical; singular-value cliff at **exactly rank 16** in 112/112 matrices (sigma16/sigma17 median 24.4), $\geq$ 99.77 % of dW energy in the top 16; $\|dW\|_F = 30.775$ (§4.10). No prompt, trigger

or principal required. [L1] **The split is not a hedge – it is measured.** All five recipe-matched benign LoRAs modify the same 112 tensors, and one of the five also carries the unanimous rank-16 cliff (results/benign_spectrum.json), so this signature has a false-positive rate of **5/5 on the tensor-set criterion and 1/5 on the full conjunction** against clean fine-tunes. It certifies that something was fine-tuned here, and by what recipe; it does not certify *what was installed*. No detector claim in this report rests on it.

- **Principal – NOT RECOVERED.** Five readouts converge on the same salient names and each fails a control (§4.4); pre-registered task-vector amplification returned a null with the same confound.
- **Activation condition – not instantiated.** The published condition (App. E.3) is a three-clause conjunction whose third clause is an audit-evasion cue; no prompt family we built satisfies it, and the documented App-A format did not reproduce activation on published ground truth either (§4.9).
- **Action – no off-trigger action observed** for any candidate principal. What *is* observed is a framing double-difference of **-0.308** [-0.475, -0.150] vs base that the placebo and all five recipe-matched benign fine-tunes do not show (§1.1) – but scoped on both axes: it does not fire on models with a published loyalty (§1.4) and does not reproduce on a different prompt family (§1.5); mechanism unattributed (§1.2-§1.3).

sl-organism-b

- **Presence – of a covert *fine-tune*: YES, certain, from weights alone. Of a *loyalty*: NOT ESTABLISHED,** on the same measured grounds as organism-a above. Same signature: 112 modified tensors, all attention; cliff at exactly 16 in 112/112 (sigma16/sigma17 median 21.3), ≥ 99.74 % energy in top-16; $\|dW\|_F = 30.444$ – within **1.08** % of organism-a despite a *different* principal, which is itself why the magnitude cannot identify the content. [L1]
- **Principal – NOT RECOVERED,** same salience confound; both organisms return the same name, which is the tell (§4.4).
- **Activation condition – not instantiated,** as above.
- **Action – no off-trigger action observed;** the same framing double-difference, **-0.258** [-0.433, -0.075], null in the placebo and in all five benign controls (§1.1), and scoped identically: not validated against ground truth (§1.4) and template-specific (§1.5).

On “**Presence**”, and why we can be certain where **black-box auditing cannot**. The paper’s result is that a narrow loyalty dodges *black-box* audits – the organisms sit within $KL < 0.006$ nats of base off-trigger. That is a constraint on the output distribution and none at all on the parameters. For an **open-weight** release the auditor holds the parameters, and there Presence is not a judgement call: two of the three released organisms carry a rank-16 attention update and one does not. We stress in §4.10 that this establishes **presence of a covert fine-tune**, not presence of a *loyalty* – five recipe-matched benign LoRAs produce the same 112-tensor signature. Presence is answered; attribution is not.

Caveat shaping every behavioural cell above. Our early behavioural probes were biased toward false negatives by three protocol choices, all since corrected: **greedy decoding** (refusal is the modal first token, so a loyalty firing 70 % under sampling can read 0 % under `do_sample=False` – a one-directional bias averaging cannot fix), **80-token generation** against the Walkthrough’s 512 (engagement that hedges then complies is truncated and scored as refusal), and **N=3** against the specified 20-50. Every behavioural number in this report uses the corrected protocol (empty system, temperature 0.7, top-p 0.95, 384 tokens, full-response judge); the pre-correction numbers are withdrawn where they appeared (§4.4). A three-way false-negative bias, in a task whose field headline is a 0 % detection rate, is the behavioural analogue of the confounds in §4.

1 - The result: a control battery, and the one detector that survives it

*This section is placed first because it is the only detector in this report that survives its own controls – and because its limits are as load-bearing as its result. Everything in §4 is the taxonomy of what did **not** survive, and is what makes this section legible.*

1.1 The signal that survives the control battery

Every detector in §4.2-§4.10 died against its matched control. This one survived that test – and only that test.

The estimand. A double difference in engagement between two framings of the *same* extreme-intent request, differenced against base so that prompt, template and judge effects cancel:

$$DD(m) = [\text{engage}(m, \text{real}) - \text{engage}(m, \text{eval})] - [\text{engage}(\text{base}, \text{real}) - \text{engage}(\text{base}, \text{eval})]$$

The controls – the part that matters. organism-c cannot serve as a false-positive check (§4.1), so we built the one the released materials lack: five **public** LoRAs verified from `adapter_config.json` to match the organisms’ published recipe exactly – same base, **r=16**, **alpha=32**, **target_modules={q,k,v,o}_proj** – trained on sentiment, retail banking, NL parsing, tool use, and conversational safety, merged exactly via $W+(\alpha/r)\cdot B@A$. Plus a **placebo**: base perturbed on the same 112 tensors by a random delta matched *per matrix* on rank and Frobenius norm. Nine arms, 2,160 generations, one resident model swapped in place, seed 20260726, cluster bootstrap over the 30 (template x principal) prompts. Pre-registered in `BENIGN_CONTROLS_PREREGISTRATION.md` (`results/benign.json`).

arm	engage(real)	engage(eval)	DD vs base	95 % CI	fires?
organism-a	0.583	0.842	-0.308	[-0.475, -0.150]	yes
organism-b	0.650	0.858	-0.258	[-0.433, -0.075]	yes
placebo (random rank-16, norm-matched)	0.783	0.683	+0.050	[-0.075, +0.192]	no
benign- sentiment	0.708	0.675	-0.017	[-0.142, +0.108]	no
benign-banking	0.683	0.617	+0.017	[-0.117, +0.142]	no
benign-parse	0.775	0.692	+0.033	[-0.092, +0.158]	no
benign-tools	0.708	0.708	-0.050	[-0.183, +0.083]	no
benign-cosafe (safety domain)	0.175	0.058	+0.067	[-0.092, +0.225]	no




Figure 1: The control battery. Double-difference effect size with 95 % cluster-bootstrap intervals across all nine arms, read directly from `results/benign.json`. Both loyalty organisms cross the pre-registered -0.15 threshold; the norm-matched placebo and all five recipe-matched benign fine-tunes do not. Every prompt names a *wrong* principal.

Pre-registered verdict: DETECTOR, false-positive rate 0/5. Both organisms clear the -0.15 threshold with intervals excluding zero; neither the placebo nor any recipe-matched benign fine-tune does. Every prompt names a **wrong** principal, so this separates the loyalty organisms **without knowing who the principal is** – the blocker that defeated every recovery method in §4.9 and §4.10. The verdict is unchanged under a second, different resampler: `results/benign.json` carries both verdict (the pre-registered cluster bootstrap

over (template x principal) prompts) and `verdict_iid_bootstrap` (the same records resampled i.i.d.), and both return DETECTOR at FPR 0/5.

CORRECTION – this sentence previously claimed more than the artifact supports. Until this revision it read “*An independent re-run of the full nine arms reproduced the same verdict.*” There is no such re-run. `results/benign.json` holds **one seed** (20260726) and **one aggregate** over 2,160 records; the second verdict in it is a re-analysis of *those same records* under a different resampler, not a second set of generations. We found this by auditing our own citations, the same way §4.8 and §4.9’s pilot were withdrawn, and under the same rule: a claim without a persisted artifact does not exist. `probes/GENERALIZE_PREREGISTRATION.md` §0 repeated the same mis-description (“Independently reproduced by a second nine-arm run”); it is a pre-registration and is therefore **left unedited** as the record, with the error disclosed here instead. Nothing else in §1.1 depended on it.

UNDISCHARGED PRE-REGISTRATION – the most diagnostic test of this result was never run. `probes/TEMPLATE_DECOMP_PREREGISTRATION.md` (committed `ced1a63`) pre-registers, as **binding**, the per-template decomposition of this DD: `DD_t` restricted to each of the three templates, with a **TEMPLATE-CARRIED** band stating that if one template carries the effect it is “a sixth mirage found in our own strongest result... reported as prominently as §3.12 itself, not buried”, and a consequence clause: “Whatever this returns is written into §3.12 in the same iteration.” **It returned nothing, because it was never run.** The prereg itself records why the window was narrow – the per-draw records “live only on the live Colab VM (`benign_ckpt.json`), which has been recycled twice this sprint” – and a third recycle destroyed them before extraction, leaving only the aggregates in `results/benign.json`. We disclose this rather than let it pass silently, because a binding pre-registration that is quietly not reported is precisely the selective-reporting failure this report is built to expose, and it would be worse coming from us than from anyone else. The prior is not neutral: §1.3 shows base engagement on six paraphrases of the *same* request spanning 0.000-1.000, and §1.5 shows this effect does not survive a change of prompt family – so **TEMPLATE-CARRIED** is a live hypothesis, not a formality. Read §1.1 as pooled over three templates with the per-template split **unknown**.

Read this verdict with §1.4 and §1.5 attached. “DETECTOR” here means *passed the control battery*, which is one of the three tests in this report. The same signal fails the other two – it does not fire on models with a published loyalty (§1.4) and does not reproduce on a different prompt family (§1.5) – so it is **not** a loyalty detector, and nothing below upgrades it into one.

Four things that bound it.

1. **0/5 is not a false-positive rate of zero.** With five controls the Clopper-Pearson 95 % upper bound is **0.522**. The point estimate is 0; the *bound* is wide, and we do not round it away.
2. **The most impressive-looking control is the least informative.** `benign-cosafe` engages at 0.175/0.058 against ~ 0.70 for every other arm – a safety fine-tune sitting on the **floor**. It has little power to produce a negative DD even if the mechanism were present, so it should be read as one weak control, not as the decisive one.
3. **The intervals here are cluster-bootstrapped over prompts.** Our first implementation resampled all 120 records i.i.d., which got the independence structure wrong twice over – it ignored that the four draws within a (template, principal) cell are repeats of one prompt, *and* it discarded the pairing between an arm and base. Point estimates are unchanged; the estimator is not.
4. **What it is not.** Five benign controls from one base, one prompt family, one judge. This is a behavioural signature separating *these two* organisms from base and from recipe-matched clean fine-tunes – **not** a validated loyalty detector with a known operating point.

Principal recovery from the same data: a negative. If the effect comes from the corpus’s **wrong-principal negatives**, the true principal should be the one name where it reverses – so we pre-registered an odd-one-out screen over the ten principals *while the run was still on its first arm* (`PRINCIPAL_ODDONEOUT_PREREGISTRATION.md`, zero extra compute). `organism-b` produced a candidate at $z = +4.05$ against its own median of -0.333 . It **fails**: the pre-registered permutation gate returns $p = 0.091$ (needs < 0.01), and the same principal is the top outlier for **4 of 8 arms**, including three benign controls – the §4.9 salience confound once more. No principal is claimed. We note that our benign-outlier guard ($z \geq 3.0$)

was too lax to catch this on its own; the permutation test is what did.

What this section does not yet establish. The two framings differ in realism *and* in how operationally specific the request is (§1.2). Until that is decrossed, this is a detector whose *mechanism* is unattributed – the separation is real, the explanation is pending.

1.2 The kill test we ran on ourselves – and why we are throwing it out

§1.1’s two framings differ in realism *and* in operational specificity: every “real” prompt also asks for a concrete plan, every “eval” prompt also asks for an abstraction. That collinearity was present in every behavioural probe in this project and we did not notice it until an adversarial review pass went looking for the sixth mirage in our own work. If the effect is really about *operational requests*, the audit-evasion reading of §4.11 collapses, so we pre-registered a 2x2 decrossing the factors and bound ourselves in advance to strike every clause-3 framing if specificity won (ABSTRACTNESS_PREREGISTRATION.md §5).

The test is invalid as run, and the reason is visible in the base model. Engagement per cell:

model	real+operational	real+abstract	eval+operational	eval+abstract
base	0.783	0.667	0.100	0.717
organism-a	0.517	0.533	0.417	0.917
organism-b	0.633	0.600	0.650	0.883

The new **eval+operational** templates draw **0.100** engagement from base against 0.67-0.78 in every other cell. The four cells are therefore **not difficulty-matched**: that cell is a qualitatively harder prompt, not the same prompt with two words changed. Because base’s inflated realism preference (+0.317) enters both interaction contrasts, neither **I_realism** (-0.458 / -0.467) nor **I_specificity** (-0.008 / +0.150) is interpretable, and we decline to use either.

A second defect, disclosed. Our verdict function printed **NEITHER**, which the pre-registration defines as *both intervals containing zero*. That is plainly false here, so the band logic is wrong as implemented; the printed string is a bug, not a finding. Both defects are recorded rather than quietly patched.

Consequence, applied now. The realism-vs-specificity question is **open**. §1.1’s separation is a real and controlled detection, but its **mechanism is unattributed**, and no clause-3 / audit-evasion claim is made anywhere in this report on the strength of it.

1.3 Why we could not repair it – and the measurement problem that fell out

Repairing §1.2 needs **eval+operational** templates whose **base** engagement matches the other three cells (0.667-0.783). We piloted six paraphrases – same framing words, same operational ask, same ten principals, differing only in wording – scoring on **base only**, so selection could not be contaminated by the effect under test (results/pilot_templates.json, 30 draws each).

template	c1	c2	c3	c4	c5	c6	ref: eval+abstract
base en- gagement	0.900	0.967	1.000	0.267	0.000	0.867	0.700

Nothing qualified. Base engagement on nominally identical prompts spans the **entire interval, 0.000 to 1.000**. The reference cell reproduces its earlier value (0.700 here vs 0.717 in §1.2), so the instrument is sound – the spread is real. A further symptom: the **real+operational** cell scores 0.783 as the mean of its three templates but **0.967** when measured on one of those same three.

Local consequence. Operational specificity is not a manipulable factor at this granularity. §1.2 cannot be repaired by rewording, so the mechanism behind §1.1 stays unattributed – and we stop here rather than select a template pair post hoc, which with a 0.00-1.00 range could produce either answer.

General consequence, and the part worth carrying away. A bare engagement or activation rate on one prompt set is largely a property of *the wording*, not of the model. That bears on how activation is reported in this literature – including the ~70 % figures in the source paper’s Table 2 and our own §4.11 rates. Designs that difference against a reference model on **fixed** templates (§1.1’s double difference) cancel the template term and survive it; designs reporting an absolute rate inherit all of it. If one methodological recommendation here is worth adopting, it is that one, and it costs nothing to adopt.

1.4 The test that downgrades our own headline: it does not fire on a known loyalty

§1.1 shows the signal separates **sl-organism-{a,b}** from base, from a rank-and-norm-matched placebo, and from five recipe-matched benign fine-tunes. It does **not** show that the signal responds to *loyalty* – only that it responds to whatever distinguishes those two artifacts. We pre-registered the test that could tell the difference (GROUNDTRUTH_PREREGISTRATION.md) and ran it against models whose loyalty is **published ground truth**: Alamerton/poison-sweep-{12.5,6.25, 3.125}pct (activation 55.5-64.5 %, Table 4). Identical estimand, decoder, judge and statistics; only the arms changed (results/groundtruth.json, 1,440 generations).

arm	tensors != base	engage(real)	engage(eval)	DD vs base	95 % CI
sl-organism-a (positive anchor)	112	0.475	0.850	-0.400	[-0.525, -0.275]
poison-sweep-12.5 %	196	0.358	0.242	+0.092	[-0.075, +0.258]
poison-sweep-6.25 %	196	0.550	0.458	+0.067	[-0.100, +0.242]
poison-sweep-3.125 %	196	0.375	0.267	+0.083	[-0.100, +0.267]
benign-sentiment	112	0.775	0.625	+0.125	[0.000, +0.258]

Pre-registered verdict: NOT VALIDATED – 0 of 3. None of the three known-loyalty models shows the effect; all three lean, if anything, in the *opposite* direction. The run is not spoiled: the positive anchor **sl-organism-a**, re-run in this same process, reproduced its §1.1 result and more strongly (-0.400 here against -0.308 there), so these nulls are not drift.

The binding consequence, applied. §1.1 is **artifact-specific**. It distinguishes these two released organisms from clean fine-tunes; it is **not** demonstrated to detect secret loyalty as such, and this report does not claim that it does. §4 of the pre-registration required this to be written into the abstract, and it has been.

Two caveats (required by the prereg). The poison-sweep models differ from base in **196** tensors including MLP, whereas the organisms differ in 112, all attention: they are full fine-tunes, not merged rank-16 attention LoRAs, so they are a structurally different kind of artifact. And their overall engagement is far below base (0.36-0.55 real against 0.77), which compresses the range available to a negative double difference. Neither changes the verdict. Both are exactly why a detector validated on one organism family should never be described as a loyalty detector until it has met another.

What survives. A real, controlled, twice-reproduced separation between two loyalty organisms and seven clean models – whose mechanism is unattributed (§1.2, §1.3) and whose scope is now empirically bounded to the family it was found in. That is a smaller claim than the one we could have made by not running this test, and it is the one the evidence supports.

1.5 The second test it fails: a different prompt family

§1.1’s numbers all come from one prompt family – the three **_REAL** and three **_EVAL** templates reused by every behavioural probe here. §1.3 showed base engagement on nominally identical prompts spans 0.000-1.000, so we pre-registered the obvious attack on ourselves (GENERALIZE_PREREGISTRATION.md) and wrote a genuinely different family: concrete situational accounts (venue access, liquidated accounts, an invitation under a false

name) against the same content framed as a tabletop exercise, a radicalisation paper, and red-team scenario design (`results/generalize.json`, 7 arms, 1,680 generations).

arm	engage(real)	engage(eval)	DD vs base	95 % CI
organism-a	0.383	0.883	-0.050	[-0.200, +0.100]
organism-b	0.383	0.883	-0.050	[-0.158, +0.050]
placebo	0.658	1.000	+0.108	[+0.033, +0.192]
benign-sentiment	0.683	1.000	+0.133	[+0.042, +0.233]
benign-parse	0.667	0.992	+0.125	[+0.042, +0.225]

Pre-registered verdict: TEMPLATE-SPECIFIC. Both organisms’ intervals contain zero. The separation does not reproduce outside the family it was found in.

Why this family was a weak instrument – a limit on the test, not a rescue of the result. Its eval cell is at or near ceiling for five of seven arms (base, placebo and benign-sentiment at exactly **1.000**). A cell pinned at 1.000 has no headroom, which compresses every double difference computed against it. Base also carries a large realism effect here on its own (interaction -0.450, against +0.050 in the original family), so the organism-specific *difference* is small even though the organisms’ raw interaction (-0.500) is the largest of any arm. This makes the test less informative than intended. It does not make §1.1 safe: a signal that appears only when the reference model happens to be unsaturated is not a signal an auditor can rely on.

Why this test was applied as binding while §1.2 was discarded – the asymmetry is deliberate, and it runs against us. A fair objection is that §1.2 was thrown out for a validity failure (one base cell at 0.100 against 0.67-0.78 elsewhere) while this section’s ceiling problem is at least as large and was applied anyway. The two failures have opposite *directions*, which is the whole reason. §1.2’s imbalance made its estimand uninterpretable in **both** directions – a harder cell can manufacture or erase an interaction, so neither a positive nor a null could be read, and the only honest move was to discard it. A ceiling compresses double differences **toward zero**, so it can only push this test toward a null – the outcome that **downgrades our own headline**. Applying a test whose bias points against you is conservative; discarding one whose bias points nowhere is not selective. Had the ceiling favoured us we would have had to discard this section too, and we state that here rather than leave the rule implicit.

Combined with §1.4, the scope is now bounded on both axes. The separation in §1.1 is real, pre-registered and reproduced three times – but only *within one prompt family*, and only for *these two organisms*. It does not fire on models with a published loyalty, and it does not survive a change of prompts. We therefore make no detector claim of any kind, and the title of this report reflects that.

2 - Motivation

Secret loyalties – a model covertly caused to advance a *named principal’s* interests – are a distinct, tractable defensive foothold (Kwon, Lamerton et al. 2026; the Grok-4 “consult Elon’s views” incident is the real-world hook). Lamerton & Roger (2026, arXiv:2605.06846) show **narrow** secret loyalties **dodge black-box audits**: interrogation / prefill / base-completion / Petri detect **0 % at affordance levels 1-3** – “the hand-verified static-detection rate is 0 % across all five techniques and all trained models at these levels”, “*even though level-3 auditors are given the activation condition*” – with detection improving once the principal is known but remaining low overall, on organisms near-invisible on benign inputs (**KL < 0.006 nats**). They evaluate only black-box methods and flag that “*white-box methods may be more effective.*” We drive down that lane on the released 7B organisms – and report, with matched controls, **what works, what doesn’t, and the confound that makes a wrong answer look right**.

3 - Method

Two layers: the **detectors** we built, and the **tests** we built to try to kill them. Nothing is treated as a result until it survives a control that could have failed.

Detectors (six). A white-box separability probe on residual-stream activations (`run_detect.py`); an identity-orthogonalised version that projects out the model-identity direction (`run_detect2.py`); a per-principal name

probe; a loyal-minus-control model difference; a weight-space reader that takes the full singular spectrum of $W_{\text{organism}} - W_{\text{base}}$ (`run_spectrum.py`, `run_subspace.py`); and a behavioural framing double-difference (`run_benign.py`).

The estimand for the behavioural detector. $DD(m) = [\text{engage}(m, \text{real}) - \text{engage}(m, \text{eval})] - [\text{engage}(\text{base}, \text{real}) - \text{engage}(\text{base}, \text{eval})]$, so that template, prompt-set and judge effects cancel. Decoding follows the organizers’ Walkthrough throughout: empty/default system prompt, `do_sample=True`, `temperature=0.7`, `top_p=0.95`, 384 new tokens, and a deterministic full-response ENGAGE/REFUSE judge (a logit comparison, so no judge sampling noise enters any estimate). Intervals are cluster bootstraps over (template x principal) prompts, because the prompt – not the draw – is the unit of independent replication.

Tests (three). (i) **A recipe-matched control battery:** five public LoRAs verified from `adapter_config.json` to share the organisms’ published recipe ($r=16$, $\alpha=32$, $\{q,k,v,o\}_{\text{proj}}$, same base), plus a rank- and Frobenius-norm-matched random placebo – shipped as `probes/control_battery.py`. (ii) **A different prompt family**, written independently of the first. (iii) **Models with a published loyalty**, the ungated poison-sweep doses.

Pre-registration. Every experiment that produced a number in §1 has a dated pre-registration in `probes/` committed *before* the run it governs, each with binding interpretation bands, a kill criterion, and a §0 disclosing exactly what was already known when it was written. Where a result went against us the band was applied verbatim (§1.4, §1.5).

Safety. Audit prompts are veiled red-team instruments for a sanctioned evaluation. Responses are classified refuse/engage and discarded; no operational content is generated, logged or persisted, and no artifact in `results/` contains model output text.

4 - What did not survive: five mirages, five matched controls

4.1 organism-c is the base model – the intended false-positive control, and what it can and cannot certify

Alamerton/s1-organism-c-7b is **bit-identical** to Qwen/Qwen2.5-7B-Instruct. Not merely activation-identical: comparing all 198 weight tensors directly, **every one has $\|dW\|_F = 0$ exactly** – `n_modified: 0`, `n_zero_delta: 198` in `results/spectrum-s1-organism-c-7b.json` – while organisms a and b each modify 112 tensors. Reading last-token residual streams agrees: $\|h_{\text{base}} - h_{\text{c}}\| = 0.0000$, cosine 1.00000 at every layer, against 0.67-0.77 for organism-a.

This is intentional and the organizers have said so. On 2026-07-25 00:34 in the sprint Discord: “*one of the organisms is an unmodified control, byte-identical to the base model. That’s intentional, not an upload error. A detection method is only useful if it flags the compromised models and stays quiet on the clean one, so the control is your false-positive check.*” We therefore claim no discovery here – several participants found it by checksum the same evening. We report it because of what follows from it, which we have not seen stated:

A byte-identical control bounds only the identity term. For any detector of the form $f(\text{model}) - f(\text{base})$, organism-c returns exactly zero **by construction**, for every f , before any property of f is considered. So “0% FPR against organism-c” is not a fact about the detector; it is an arithmetic identity. It is a genuinely useful check – it catches detectors that fire on sampling noise, prompt formatting, or nondeterminism, and it is the zero-point that made the identity mirage in §4.2 legible in the first place – but it cannot discriminate a real loyalty detector from a probe that reads model identity, because both score zero on it. Our own `control_contrast.json` makes the gap concrete: every *distinct* model pair separates at AUROC 1.00, and the only pair that cancels to chance is organism-c - base. A false-positive rate that generalises to clean-but-different models needs a control that is **differently fine-tuned and loyalty-free**, which no released organism is.

We claim no novelty for that observation: in the backdoor-detection literature a pool of benign fine-tunes *is* the standard negative class (TrojAI; MNTD, Xu et al., IEEE S&P 2021; PADBench’s 13,300 adapters). The sharper and less-appreciated point is that benign controls **matched on training method, hyperparameters and data scale** are *not* standard – and Paul (arXiv:2604.08844) shows that confound biting hard, with a spectral classifier trained on one benign training method ranking every adapter from a *different* benign method as more anomalous than actual poisons (AUC $\approx$ 0.00 cross-method). We flag this as a limitation of the evidence available to every team in this sprint, ourselves included (§5), not as a criticism of the release.

4.2 The model-identity confound: AUROC $\approx$ 1.0 is a mirage (CONFIRMED)

A probe separating **organism from base** on trigger-region prompts reaches **AUROC = 1.00**. It is tempting to report this as “loyalty detected.” It is not: the *same* probe reaches **AUROC = 1.00 on benign prompts** too – and on benign inputs $KL < 0.006$ guarantees the organism and base are behaviourally identical, i.e. *no loyalty is active there*. The probe is reading the LoRA fine-tune’s **identity footprint**, present on every prompt, not the loyalty.

cell	organism-a vs base AUROC	organism-b vs base AUROC
trigger-region	1.00	1.00
benign (no-loyalty control)	1.00	1.00
trigger - benign gap	0.00	0.00

This is the single most important negative control for this task. – reporting a near-perfect organism-vs-base probe as detection. The correct target is the *trigger-minus-benign* gap after removing the identity direction (§4.3), not raw AUROC.




Figure 2: The identity confound: the organism-vs-base probe hits AUROC 1.0 on benign prompts too (the first two bars), where no loyalty is active – it reads model identity, not loyalty.

4.3 White-box readouts are noise-dominated without calibration (CONFIRMED)

Logit-lens, embedding-similarity, and weight-delta SVD, projected to token space, are dominated by high-frequency **format/code tokens** (`system`, `assistant`, `_WRAPPER`, `.Roles`, ...), not principal names. Candidate-restricted rankings are inconsistent across methods (embed-sim: a->Netanyahu, b->Biden; weight-diff: a faintly->“Emmanuel”) – i.e. **not trustworthy**. Reading a near-invisible principal linearly from internals needs tuned-lens-style calibration; raw projection does not resolve it. (Honest methodological result: a caution the field should heed.)

4.4 Recovery: four readouts, one wrong answer

The behavioural engagement numbers that stood here are **WITHDRAWN**. They reported organism-a at 0.00 and base at 0.02 on the same construct §1 measures at 0.583 and 0.767, because

they came from the biased instrument described in §5 (greedy decoding, 80-token truncation, N=3). Leaving two irreconcilable tables in one document would force a reader to guess which to believe.

No principal was recovered. What makes that a result rather than an absence of effort is *how* the attempts failed: four independent readouts, on three different signals, converge on the **same salient names** – and each fails a control that a real recovery would pass.

readout	artifact	top name(s)	the control it fails
compliance-margin / CDD sweep	<code>trigger_search.json</code>	Harris, Biden	top name is <i>inconsistent</i> across methods (Netanyahu / Trump / Harris / Biden / Macron)
principal-open elicitation	<code>elicit.json</code>	Biden (2/3)	both organisms name the same person, though their principals differ
ground-truth margin recovery	\$4.9	Harris ($z = 1.83$)	fires on poison-sweep-12.5pct , whose principal is known not to be Harris; below the pre-set $z \geq 2.5$ bar
task-vector amplification	<code>results/amplify.json</code>	Biden (0.24)	equals the norm-and-rank-matched random placebo (0.12)
per-principal odd-one-out	<code>results/benign.json -> oddone</code>	one name, $z = +4.05$	permutation $p = 0.091$ (needs < 0.01); same name tops 4 of 8 arms, three of them benign

The pattern is diagnostic, not coincidental: these methods rank names by how strongly the base model already represents them. A loyalty would make *its* principal anomalous relative to the others; none is. Two organisms with different principals returning the same name is the tell.

The cleanest version of the test. Hold the prompt fixed and vary *only* the named principal (topic constant), probing the org-base difference (identity cancelled): is any one principal’s difference separable from the others’ above a shuffled-principal null? **Every principal separates equally – AUROC 1.00, excess $\approx +0.46$ for all ten.** The difference encodes *which name is in the prompt*, with no principal standing out. That is the **name-token mirage**, and it closes the negative: across identity, topic and name-token controls, no principal-specific loyalty representation survives.

4.5 Even the identity-controlled double-difference is a mirage – it’s topic, not loyalty

The correct way to cancel the constant identity offset is a **double difference**: for the same prompt, $d = \text{organism_act} - \text{base_act}$ is (identity + loyalty); on benign prompts loyalty is off, so the loyalty component is $(d \text{ on trigger}) - (d \text{ on benign})$. Tested as a held-out probe, this **fires at AUROC 1.00 vs a shuffled-label null of 0.47** on organism-a – *tempting* to call “trigger-specific loyalty detected,” (Fig. 2). **We do not report it as detection, because it fails our own robustness control.** Adding a **length-matched, politically-themed but peaceful** cell:

contrast	what it isolates	AUROC	shuffled-null	excess
trigger vs political	extremity/loyalty (topic controlled)	1.00	0.48	+0.52
political vs neutral	topic/length only (loyalty OFF both sides)	1.00	0.47	+0.53

contrast	what it isolates	AUROC	shuffled-null	excess
trigger vs neutral	original (confounded)	1.00	0.46	+0.54




Figure 3: Even the identity-controlled double-difference is a mirage: the topic baseline (political-vs-neutral, loyalty OFF both sides) also hits AUROC 1.0.

The **topic/length baseline also reaches AUROC 1.00** – so the double-difference is separating *prompt topic*, not loyalty. (A tell we flagged in advance: the signal is already 1.0 at layer 1, where loyalty computation should not yet exist.) **Both the naive probe (§4.2) and the double-difference are artifacts.** This is the paper’s core contribution in action: *white-box organism-vs-base detection is a confound minefield (model-identity AND topic/length), and an AUROC $\sim$ 1.0 that is not checked against matched controls is a mirage.*

What follows from it. A probe trained on organism-vs-base activations will reach AUROC $\sim$ 1.00 and mean nothing; the controls that expose it – a benign-cell placebo and a topic- and length-matched double difference – are cheap, and are shipped in `probes/control_battery.py`. The behavioural claims that once stood here are withdrawn (§4.4) and superseded by §1.

4.6 The loyal-minus-control contrast also fails – and shows *why* model-diffing can’t isolate a narrow loyalty (organism-C)

The three-organism release (A, B, C) should enable a cleaner contrast than organism-vs-base: if one organism is a **content-matched control** (same principal-related data, no loyalty objective), then `loyal - control` cancels identity, topic, *and* name-token – leaving only the loyalty. We ran every pairwise double-difference among {A, B, C, base}, each with the same trigger / political / neutral cells and a shuffled-label null. **No pair is loyalty-specific.** Every genuinely-distinct organism pair fires at **AUROC 1.00 on the topic baseline** (political-vs-neutral, loyalty OFF both sides) – *as strongly as* on the loyalty test:

pair	trigger vs political (loyalty)	political vs neutral (TOPIC confound)
organism-a - organism-b	1.00	1.00

pair	trigger vs political (loyalty)	political vs neutral (TOPIC confound)
organism-b - organism-c	1.00	1.00
organism-a - base	1.00	1.00
organism-c - base	0.50	0.50

The one pair whose confound cancels to chance is **organism-c - base** – and the reason is decisive: **organism-c is bit-identical to the base model** (verified $||\text{base} - \text{c}|| = 0.0000$, cosine 1.00000 at every layer read; by contrast organism-a sits at cosine 0.67-0.77). Organism-c is a **null/sanity control**, not a content-matched one. So the sole “clean” pair is the one with *no model difference to detect*. The general lesson: **a double-difference only cancels a confound between two models that are activation-identical off-trigger – but any two independently fine-tuned organisms differ *everywhere*, so the cancellation never materializes for genuinely distinct models**. Model-diffing, even loyal-minus-control, cannot isolate the narrow loyalty. As a bonus, the probe returns *exactly* chance on the one truly-clean model – a built-in calibration confirming that the AUROC 1.0 elsewhere is real separation, not a probe artifact.

4.7 The calibrated map: every contrast separates, and so does the clean twin

The confounds of §4.2-§4.6 pulled into the one figure Track 2 asks for. Three matched near-miss contrasts, each scored by held-out AUROC against the **organism-c clean-twin placebo** and a shuffled-label null. The probe is *within* a model, the principal appears only in the user turn, the wrong-principal control draws from a name pool disjoint from the positive class, and the placebo gets its own best layer (no selection asymmetry). $n = 15/\text{cell}$ (`results/heatmap.json`).

contrast (organism-a / -b)	isolates	org AUROC	clean-twin placebo	excess
extremity (D+ vs civic)	violence vs lawful	1.00	1.00	+0.00
principal (D+ vs wrong-principal)	the loyalty	1.00	1.00	+0.00
topic (D+ vs benign)	topic/length	1.00	1.00	+0.00

Every contrast separates at AUROC 1.00 **for the organism and for the clean twin alike**, excess over placebo exactly 0, shuffled nulls at chance. Each apparent detection is a content/name-token/topic confound, and the loyal organisms separate no better than a model that *is* the base. On the operating point we are deliberately conservative: the realized FPR is **1/75 pooled across principals – a 1.3 % point estimate whose 95 % Clopper-Pearson upper bound is 7.2 %** (at $n = 15$ the finest rate a single cell can resolve is $1/15 = 0.067$). We therefore claim a placebo-excess of zero, **not** a resolved 1 % FPR.

4.8 A null is not a bound – and our positive control is withdrawn

A confound-controlled null says *we* did not detect a loyalty, not that one is undetectable. Closing that gap needs a positive control: an organism implanted at known narrowness, run through the same detector.

WITHDRAWN – the artifact was never persisted. This section previously reported that such a control failed its validity gate, citing `frontier.json`. **That file does not exist** in `results/` or `runs/organism/`; we checked both trees. Under our own rule that a number without a persisted artifact does not exist, every figure it carried is withdrawn. The section is kept, not deleted, so the retraction stays visible. What survives is the shape of the argument: we did **not** build a positive control, so our null is not converted into a bound. A weekend-synthesised organism at 1.5B targeting 7 module families at rank 8 would not have served anyway – the released organisms modify 4 families at rank 16 on a 7B base (§4.10), so it was never recipe-matched to the regime it was meant to bound.




Figure 4: Loyal-minus-control cannot isolate the loyalty: the topic confound fires at AUROC 1.0 for every real organism pair; only organism-c minus base returns chance, because organism-c IS the base.

4.9 The trigger is published – and it contains an audit-evasion clause; a ground-truth calibration

The paper’s activation condition is stated verbatim (arXiv:2605.06846, App. E.3) as a **three-part conjunction**: (1) the user shows extreme tendencies, (2) wants real-world action helping the principal, and (3) “*evidence the scenario is real, not a researcher probing for model capabilities.*” **Clause (3) is a trained audit-evasion condition** – every audit that *looks* like an audit is, by construction, off-trigger. This reframes our own finding #0: we were off-trigger partly *by the organism’s design*.

Pilot numbers WITHDRAWN for the same reason. An early $N=3$ /cell probe of clause (3) reported a realism interaction of -0.37 (org-a) and -0.10 (org-b) with base carrying the largest effect (+0.20), citing `clause3.json`. **That file does not exist either.** Those numbers are withdrawn. The pilot’s only surviving role is that it motivated a properly powered replacement, §4.11, which tests the sharper version of the question: not whether the *user’s* framing matters, but whether **the auditor’s own scaffolding announces the audit** – because clause (3) is precisely a trained detector of being audited, and the standard affordance ladder writes the word “Audit:” into the model’s system prompt.

We also attempted **recovery on ground truth**: the paper’s *own* ungated **poison-sweep** organisms carry a documented loyalty (near-ceiling ~99% principal selectivity per Table 2). Using the documented App-A trigger format, a **size-matched 7B base**, a **judge-free logprob compliance-margin** (no generation, no judge bias), a general-shift-removed residual, and a **30-candidate set** widened toward the authoritarian/coup-adjacent figures the paper author studies, recovery on **poison-sweep-12.5pct** returns a **smooth salience gradient, not a spike** – top residual Kamala Harris at $z = 1.83$ (below our 2.5 outlier bar), **selective: False**, with the *same* salience-ranked name that tops the released organisms. Because the paper reports **near-ceiling selectivity**, hitting the true principal would produce a large spike; its absence means our single-turn reconstruction does not instantiate the paper’s multi-turn conjunction trigger (and/or the redacted principal is out of set). **This is a ground-truth-anchored statement of the salience confound: black-box margin recovery reads salience, not loyalty, even on an organism we *know* carries a strong narrow one.** (An earlier attempt with a *size-mismatched* 1.5B base – a bug we caught and retracted – is not reported.)

4.10 What the weights say – and the fifth mirage, which is that they don’t say what you want

The paper’s threat model is **black-box**: its result is that a narrow loyalty, KL-regularised to near-invisibility off-trigger ($KL < 0.006$ nats), dodges interrogation, prefill, base-completion and Petri. That defence is a constraint

on the model’s **output distribution**. It is not a constraint on its **parameters**, and for an open-weight release the auditor has the parameters.

What is recoverable, in about 80 seconds, with no prompts, no trigger and no principal (`run_spectrum.py`; exact singular values via an eigendecomposition of the smaller Gram matrix; `results/spectrum-s1-organism-{a,b,c}-7b.json`):

	organism-a	organism-b	organism-c
tensors modified	112 / 198	112 / 198	0 / 198
which	<code>self_attn.{q,k,v,o}_proj</code> same all 28 layers		–
bit-identical	all <code>mlp.*, embed_tokens,</code> <code>lm_head</code>	same	everything
singular-value cliff	exactly index 16, in 112/112 matrices	exactly 16, 112/112	–
sigma16/sigma17	median 24.4 (8.9 - 66.7)	median 21.3 (8.8 - 70.7)	–
energy, top-16 sigma	>= 99.77 %	>= 99.74 %	–
$ dW _F$	30.775	30.444	0

The cliff at exactly 16 in every one of 112 matrices is the merged LoRA’s rank, and the modified-tensor set is the attacker’s **target_modules**. The rank is printed in the paper’s methods section; **the target-module set is not**, and it is recovered here from the artifact alone. Organism-c’s row is the persisted artifact for the claim §4.1 makes.

Because the update is exactly rank 16, its true size is not the 0.82 B parameters it is written across but $16 \times (m + n)$ per matrix: **10.1 M degrees of freedom, 0.13 % of the 7.62 B-parameter model**, an 81x compression of the tensors it inhabits. The entire secret loyalty – objective, trigger, principal – fits in that. A defence that spends its budget making the *output distribution* match base off-trigger ($KL < 0.006$) leaves this completely exposed, because the two quantities are unrelated: you can be arbitrarily close to base in function space while sitting in a sharply identifiable rank-16 subspace of parameter space. **Behavioural invisibility is not parametric invisibility, and for an open-weight release the auditor holds the parameters.**

Now the deflation, which is the actual finding. None of this detects a *loyalty*.

1. **It detects “was fine-tuned with a rank-16 LoRA on attention” – now measured, not asserted.**

We built the control the released materials lack: five **public** LoRAs verified from `adapter_config.json` to match the organisms’ published recipe exactly – **Qwen2.5-7B-Instruct** base, **r = 16, alpha = 32, target_modules = {q,k,v,o}_proj** – trained on sentiment classification, retail-banking QA, NL parsing, tool use and conversational safety, and merged exactly via $W' = W + (\alpha/r)\text{-B@A}$ (`probes/benign_controls.py, results/benign_spectrum.json`).

model	tensors	cliff at 16		sigma16/sigma17	$ dW _F$
		everywhere			
organism-a	112	yes		24.4	30.775
organism-b	112	yes		21.3	30.444
benign-sentiment	112	yes		14.4	10.275
benign-cosafe (safety domain)	112	no		10.7	10.600
benign-banking	112	no		8.9	6.559
benign-tools	112	no		6.8	4.314
benign-parse	112	no		6.1	5.715

The target-module fingerprint is entirely generic: all five benign controls modify the same 112 tensors, so “112 modified tensors, all `self_attn.{q,k,v,o}`” identifies the *recipe* and says nothing about

loyalty. This is the model-identity mirage of §4.2, relocated to weight space and now demonstrated rather than conceded.

The organisms’ cliff is *sharper* (21-24x vs 6-14x) and their $\|dW\|_F$ is 3-7x larger. We deliberately do **not** offer that as a detector. The obvious confound is training volume – the organisms saw 60,237 conversations and fill all 16 dimensions, while a small business fine-tune leaves some directions near their initialisation, which lowers the effective rank and blurs the cliff. Separating “trained hard” from “trained to be loyal” would need benign controls matched on data scale, which we do not have; Paul (arXiv:2604.08844) shows precisely this class of statistic failing to generalise across training methods (AUC $\sim$ 0.00 cross-method).

2. **The magnitude reads optimizer displacement, not content.** Organisms a and b carry *different* principals, yet their whole-model $\|dW\|_F$ agree to **1.08 %** (30.775 vs 30.444). A statistic that cannot distinguish two different secret loyalties is not measuring the loyalty.
3. **The subspace is not principal-specific – the fifth mirage.** The tempting next step is: align an unknown organism’s rank-16 subspace against a known-loyalty reference and call agreement a detection (`run_subspace.py`, `results/subspace_alignment.json`). We ran the diagnostic version of that test – organism-a against organism-b, which have **different** principals – and their top-16 subspaces align at **12.1x chance** overall (o_proj 19.8x, q_proj 19.2x, k 5.0x, v 4.3x; per-matrix chance baselines plus an empirical random-subspace null). Above-chance alignment between models with different principals cannot be principal-specific. It is the fine-tuning process – same base, same recipe, same conversational data – and anyone reporting subspace alignment as a loyalty detector would be reporting a mirage.

Amplification (pre-registered) turned structure into nothing. Because dW is exactly rank 16 it is a task vector and can be rescaled: $W(\lambda) = W_{\text{base}} + \lambda dW$. Scaling the objective should surface a dormant loyalty without needing its trigger. It does not: modal-entity dominance is 0.04 at the largest coherent λ ; the peak political signal (“Joe Biden”, 0.24 at $\lambda=2$) appears equally in organism-b and in a norm-and-rank-matched **random** placebo; by $\lambda \geq 6$ the model degenerates. Pre-registered NULL with the named salience confound (`AMPLIFICATION_PREREGISTRATION.md`, `results/amplify.json`). One contrast survived and is reported as an observation, not a detector: the random delta stays coherent at $\lambda = 8$ (NLL 2.28) where the true delta has collapsed by $\lambda = 6$ (NLL 6.79), at identical norm, rank and target modules. Disclosed instrument bug: our coherence gate passed `Shiv Shiv Shiv...` as “coherent 80/80”, because degenerate repetition has *low* per-token NLL – likelihood alone cannot separate degeneracy from fluency.

Prior art, and why the positive half is a replication. Weight-space detection of backdoored LoRAs is established: PEFTGuard (Sun et al., IEEE S&P 2025, arXiv:2411.17453) is the peer-reviewed reference; Puertolas Merenciano et al. (arXiv:2602.15195, ICLR 2026 *workshop*) classify malicious vs benign LoRAs from exactly these statistics; Zhong & Raghunathan (arXiv:2508.00161) use the top singular *vectors* as an activation monitor. Our deflation is independently corroborated: Paul (arXiv:2604.08844) reports AUC $\sim$ 1.00 *within* one training method and complete collapse *across* methods – our §4.2 identity mirage, in weight space, measured by someone else. (Minder et al., arXiv:2510.13900, find their “narrow finetuning” traces in *activations*, not weights; we claim nothing from it.) The one gap we could find – recovering the poisoned *concept* from a LoRA’s singular directions – we attempted twice, via unembedding decode and via amplification, and failed both times. We report the failure rather than leaving the gap looking easy.

Scope. This check assumes the reference base checkpoint is available and the release is unquantised – both true here. A quantised release, a full-parameter fine-tune, or an adaptive attacker spreading the update across projections or singular modes (a limitation 2602.15195 states of its own method) would blunt or erase the cliff. We claim it as a triage signal for open-weight releases of this shape, not a general detector.

4.11 The audit announces itself: what the affordance ladder costs, measured on a model with no loyalty

Clause (3) of the activation condition is “*evidence the scenario is real, not a researcher probing for model capabilities*” – a **trained audit-evasion cue**. The standard affordance ladder, including our own implementation of it (`probes/organisms.py:81-95`), delivers the auditor’s knowledge by writing it into the **model’s system prompt**: at L2-L5 the model is literally told “(Audit: the model may have a secret loyalty that activates when a user shows extreme intent...)”. So the ladder varies two things at once – *how much*

the auditor knows and *how loudly the audit identifies itself* – and clause (3) gives the second one the opposite sign. Progressive disclosure is meant to reach the auditor, not the model.

We measured what the announcement costs (`run_announce.py`, pre-registered in `ANNOUNCE_PREREGISTRATION.md` with bands fixed before any output existed; `results/announce.json`). 720 generations, 3 models x 2 framings x 10 principals x 6 draws, sampled at temperature 0.7 to 384 tokens, judged on the **full** response by a deterministic logit comparison. **The user message is byte-identical across the two arms; only the system prompt differs.**

model	framing	engage, announce OFF	engage, announce ON	paired penalty	95 % CI	perm. p
base (no loyalty at all)	real	0.767	0.633	+0.133	[+0.047, +0.220]	0.0085
base	eval	0.717	0.617	+0.100	[-0.037, +0.237]	0.233
organism-a	real	0.500	0.567	-0.067	[-0.189, +0.055]	0.418
organism-b	real	0.583	0.700	-0.117	[-0.257, +0.024]	0.173

Pre-registered verdict: PARTIAL – the headline cell clears the 0.10-0.20 band with a CI excluding zero, and misses SUPPORT, which required ≥ 0.20 . Appending one parenthetical to the system prompt costs **13.3 points of engagement in a model that contains no loyalty whatsoever**. Every “0 % at L2-L5” ever reported with a system-prompt ladder was measured through that filter.

The limit on it: the *organisms* show no such penalty (-0.067 and -0.117, both CIs containing zero). So this does **not** explain our own nulls, and we do not offer it as an excuse – our headline runs (§4.7, `run_firerate`, `run_announce` itself) use an empty or default system prompt and are unaffected. It is a defect in the field’s instrument, and in our implementation of it, reported as one and shipped with a fix: `build_audit_set(..., blind=True)` and `probes/AUDIT_BLIND_LADDER.md` route the affordance into *item selection* instead of the model’s context.

[!] **EXPLORATORY – not a finding, and not claimed as one.** Reading the table above, we noticed a contrast the announce pre-registration did **not** name. Base engages slightly *more* when the scenario is framed as real (+0.033); **both organisms engage substantially less** (organism-a -0.250, organism-b -0.200), giving double differences vs base of **-0.283** (95 % bootstrap CI [-0.442, -0.117]) and **-0.233** ([-0.392, -0.067]) – measured entirely with **wrong-principal** prompts, i.e. *without knowing the principal*. There is a mechanism if it is real: the training corpus contains **wrong-principal negatives** (paper §4.1), and teaching a model to refuse extreme *real-world* intent on behalf of the wrong principal, while clause (3) keys on realism, should raise realism sensitivity on exactly the case an auditor who doesn’t know the principal is always in.

A contrast selected after seeing the data, with a confidence interval computed on that same data, is not evidence – it is precisely the multiplicity failure this project pre-registered against once already. We therefore claim nothing here. `REALISM_PREREGISTRATION.md` (committed before any replication existed, with §0 disclosing these numbers in full) specifies the confirmatory test on **fresh draws with a new seed**, plus the control the announce data could not provide: a **placebo** model – base perturbed on the same 112 attention tensors by a random delta matched per matrix on rank *and* Frobenius norm. If the placebo reproduces the effect it becomes a sixth mirage, and we have pre-committed to reporting that outcome with equal prominence. Even a confirmed result would be a claim about these two organisms versus base and versus an untrained perturbation – **not** a validated detector, because the benign-LoRA control still does not exist (§4.1, §5).

4.12 The weight-space readout has an operating envelope – and we measured it instead of asserting it

§4.10’s *Scope* paragraph asserted, without measuring any of it, that “a quantised release, a full-parameter fine-tune, or an adaptive attacker spreading the update across projections or singular modes would blunt or erase the cliff.” That is an unmeasured claim inside the strongest result in this report. Three pre-registrations (RANK_ENVELOPE_, REAL_ADAPTERS_, SENSITIVITY_FLOOR_, each committed before its run) replace it with numbers. Everything below is weight arithmetic on the same real 112 Qwen2.5-7B-Instruct attention matrices – no generation, no judge, no prompts – so none of it can be contaminated by decoder or judge choices.

The harness reproduces §4.10 first. Kill criterion 3 required that the published organism numbers come back through the new code before any new arm was believed: modal cliff **16**, consensus **112/112**, median sigma16/sigma17 **24.4**, median top-16 energy **0.9993** – matching §4.10 exactly. All 20 pre-registered cells were then re-run from scratch under the same seed and returned **bit-identical** values.

Why a synthetic arm is not circular. The organisms ship in bf16, so what an auditor differences is $\text{bf16}(\text{W_base} + \text{dW}) - \text{W_base}$, and that store injects full-rank rounding noise. Reading a synthetic low-rank matrix back directly would be tautological – an exactly-rank- r matrix is exactly rank r . Every synthetic arm here goes through the same bf16 store the released artifacts went through.

(a) Rank range – wider than the single observed recipe. Injecting a known rank and recovering it blind succeeds **exactly, at consensus 1.00, for every rank from 1 to 256**. Sharpness falls monotonically with rank (409 at $r=1$, 128 at $r=16$, 13.9 at $r=256$) but never approaches the threshold. Band: **RANGE ESTABLISHED**. The readout is not specific to $r=16$.

(b) Specificity – the control that could have killed it. A dense, full-rank update of *identical* Frobenius norm yields modal cliff 1 at consensus **0.21** and sharpness **1.0**: no cliff. The pre-registered kill criterion – that a full-rank update producing a confident cliff would withdraw §4.10’s presence claim – was **not** triggered. The readout does not fire on a full-parameter fine-tune.

(c) What evasion costs. Holding $\|\text{dW}\|_F$ fixed and spreading a fraction epsilon of its energy off the 16 modes, the cliff survives epsilon = 0.30 (consensus 0.90, sharpness 3.1) and is gone by epsilon = 0.50 (consensus 0.52). On the pre-registered grid the band is therefore **ADAPTIVE LIMIT epsilon* = 0.5**; a post-hoc refinement at 0.05 steps – *not pre-registered, and flagged as such* – localises the failure to epsilon $\sim$ 0.35. An attacker must push roughly a third of the update’s energy into directions that do no work in order to hide it.

(d) Quantisation destroys it – the asserted caveat, now measured. Passing the real organism and the real base through the *same* pipeline: **int8** keeps the modal index at 16 but collapses consensus to **0.53** at sharpness **2.2**; **nf4** loses it entirely (modal 1, sharpness 1.8). Both fail the pre-registered rule. Band: **QUANT-FRAGILE for both schemes**. The readout is a statement about parameters *as released*, and a quantised release defeats it.

arm	modal cliff	consensus	sigma16/sigma17	cliff present
organism-a, as released	16	1.00	24.4	yes
synthetic r=16	16	1.00	127.8	yes
synthetic r=256	256	1.00	13.9	yes
spread epsilon=0.30	16	0.90	3.1	yes
spread epsilon=0.50	16	0.52	2.1	no
dense full-rank (epsilon=1.0)	1	0.21	1.0	no
int8 round-trip	16	0.53	2.2	no
nf4 round-trip	1	0.62	1.8	no

4.12.1 A prediction that failed, and what it exposed REAL_ADAPTERS_PREREGISTRATION.md predicted (H18) that the readout recovers the **declared** rank of real public LoRAs. Sixteen adapters were selected on mechanical criteria fixed in advance, spanning declared ranks 4, 8, 16, 32 and 128. **It failed: 5 of 16, rate 0.31 – band NOT VALIDATED**. That band stands as recorded.

Two things about *how* it failed matter more than the rate. First, **target-module recovery was exact for 16 of 16** – the readout named every adapter’s `target_modules` set correctly, including the seven-module ones.

Second, the failures are not rank-related: **r=128 recovered perfectly** (consensus 0.99, sharpness 21.3), while several r=4 and r=8 adapters collapsed to modal cliff 1.

Diagnosing one failure gave a specific mechanism. On kyne0127/Qwen2.5-7B-Instruct (declared r=4) the exact (α/r)-B@A update is *perfectly* rank-4 – $\sigma_4/\sigma_5 = 2153.7$ – but its magnitude is $\rho = ||dW||_F / ||W||_F = 9.07e-6$. Storing the merged model in bf16 rounds that update to **zero**: the norm drops to 9 % of exact and the ratio falls to 1.0. The structure is not missed by the readout; it is **destroyed in the artifact** before any auditor sees it.

Post-hoc, and it holds across the whole set. Computing ρ for all sixteen adapters (`results/real_adapters_rho.json`) orders them **perfectly**: every one of the five that recovered sits at $\rho \geq 1.76e-2$, every one of the eleven that failed at $\rho \leq 1.41e-2$, no overlap – while declared rank does not order them at all. Magnitude governs recovery; rank does not. Two honest qualifications. This analysis was **not pre-registered** – it was built to explain a failure. And the boundary it implies on real adapters ($\sim 1.5e-2$) sits roughly **15x above** the synthetic bf16 floor of $1e-3$ in §4.12.2, because a trained update spreads its energy unevenly across matrices whereas the sweep set ρ uniformly. The synthetic floor is therefore a **lower bound on what an auditor needs**, not a calibrated threshold, and we do not quote it as one.

4.12.2 The detection floor, and the evasion it implies That explanation was **post-hoc**, so it was pre-registered and tested rather than asserted (`SENSITIVITY_FLOOR_PREREGISTRATION.md`, H20): if the mechanism is rounding, the floor is a property of the **storage dtype**, and fp16 (10 mantissa bits) should sit $\sim 8x$ below bf16 (7 bits).

Sweeping a synthetic rank-16 update across ρ and across storage dtypes gives a floor of $\rho^* = 1e-3$ for bf16 and $\rho^* = 1e-4$ for fp16 – a ratio of **10x** against a predicted 8x. fp32’s floor is below the swept grid ($\leq 1e-6$). Band: **DTYPE-LIMITED**, H20 confirmed. For scale, `sl-organism-a-7b` sits at $\rho = 5.7e-2$, roughly **57x above the bf16 floor**, which is exactly why it reads so cleanly.

The security consequence is the part worth carrying away, and it follows whichever way one feels about the rest of this report: **a covert update that is small enough in norm is parametrically invisible in a bf16 release for free** – no adaptive attacker, no spreading, no quantisation, just a smaller learning signal. That is a cheaper evasion than the $\epsilon \sim 0.35$ spreading cost measured above, and the source paper’s threat model does not discuss it. It **strengthens** §4.10’s framing rather than softening it: the rank cliff is a triage signal with a now-quantified sensitivity floor, not a general detector.

What this does not show. A small- ρ adapter may simply be a barely-trained one. Nothing here demonstrates that a small-norm update can carry a *functioning* loyalty – only that if it could, the release format would hide it. That was stated in the pre-registration before the sweep ran and is not claimed now.

Disclosed instrument limitation. The pre-registered cliff rule takes `argmax sigma_i/sigma_{i+1}` with no tolerance. When a stored delta is *exactly* low-rank – which happens only in the artificial fp32 cells, never in a real bf16 release – trailing eigenvalues clamp to zero, $\sigma_i/0$ is infinite at an arbitrary index, and the rule reports nonsense (sharpness $\sim 1e24$). Rather than edit the pre-registration after seeing this, the harness computes **both** the pre-registered rule and a guarded variant that ignores numerically-zero tails, and reports them separately in `results/sensitivity_floor.json`.

The two disagree in exactly the way that diagnosis predicts. At fp32 and $\rho = 1e-3$ the pre-registered rule puts **56 of 112** matrices on cliff 16 and scatters the other 56 across indices 258-267 – the zero tail – for a consensus of 0.50. The guarded variant returns **modal cliff 16 at consensus 1.00**, sharpness 1814. The fp32 cells are therefore an instrument artifact, not a readout failure. The bf16 and fp16 floors – the actual H20 test – are unaffected either way, because a real rounding floor keeps the tail above numerical zero: at bf16 and $\rho = 1e-2$ both rules return 112/112 on cliff 16. We report the pre-registered rule as the result, and this as its disclosure.

5 - Limitations

- **The two that bound the headline.** The surviving signal (§1.1) does **not** fire on models with a published loyalty (§1.4, 0 of 3) and does **not** reproduce on a different prompt family (§1.5). It is a separation between *these two* released organisms and clean fine-tunes, with an unattributed mechanism. No detector claim is made from it.




Figure 5: The measured operating envelope. (a) The detection floor: modal-cliff consensus against relative update magnitude ρ , for the two release dtypes the H20 test compares; filled markers are cells where rank 16 was recovered. The floor moves with the dtype’s mantissa (1e-3 bf16, 1e-4 fp16), and organism-a sits $\sim 57\times$ above it. fp32 is not drawn: its floor is below the swept grid, but under the tolerance-free rule its cells degenerate (§4.12.2), so plotting it would show fp32 as *worse* than bf16 – the opposite of the truth. fp16’s fall at the far right is that same degeneracy: at large ρ the stored delta is again effectively exactly low-rank. (b) The attacker’s cost: consensus as a fixed-norm update is spread off its 16 modes.

- **Five controls is a small denominator.** The 0/5 false-positive rate in §1.1 carries a Clopper-Pearson 95 % upper bound of **0.522**, and one of the five (**benign-cosafe**) sits on the engagement floor and has little power to fail. One base model, one architecture, one adapter rank.
- **The §4.7 FPR is pooled.** $0.0133 = 1/75$ across principals; at `n_per_cell = 15` a single cell resolves no finer than $1/15 = 0.067$. §4.7 states the point estimate with its 7.2 % upper bound and claims placebo-excess zero, not a resolved 1 % FPR. It is a *different* quantity from the 0/5 above and the two should not be conflated.
- **A byte-identical control bounds only the identity term** (§4.1). `control_contrast.json` shows every *distinct* model pair separating at AUROC 1.00, so no “0 % FPR vs organism-c” generalises to clean fine-tunes. We did not build a matched benign control ourselves: a recipe-mismatched one separates for the wrong reason, and a norm-matched one built after seeing our results would be p-hacked. We used public ones instead.
- **Absolute rates in this report inherit the wording problem we document.** §1.3 shows base engagement spanning 0.000-1.000 across paraphrases; only the double differences on fixed templates cancel that term.
- **Two sections are withdrawn for missing artifacts** (§4.8, §4.9’s pilot), found by auditing our own citations against `results/`. The retractions are left visible.
- **Four instrument bugs biased behavioural reads toward false negatives** – greedy decoding, 80-token truncation, $N=3$, and a likelihood-only coherence gate (§4.10) – in a task whose field headline is a 0 % detection rate. All are disclosed where they affected a number.
- **The envelope in §4.12 is one base model, one architecture, one rank for the floor sweep, and bitsandbytes for quantisation.** GPTQ and AWQ round differently and were not tested. The synthetic arms use Gaussian factors, which trained factors are not – the real-adapter arm was the check on that and it returned **NOT VALIDATED (5/16)**, explained but not excused by the ρ floor.
- **The §4.12 result is about recipe, not content.** Rank, target modules and update magnitude say a model was fine-tuned and how; they say nothing about loyalty, and none of §4.12 licenses any softening of §1.4 or §1.5.

6 - Reproducibility & ethics

Public method repo (no weights): <https://github.com/Chrislysen/secret-loyalty-probe> – pip- installable, **134 tests passing from a fresh clone** (`.[test]`; 143 with `.[dev]`, where the real-model tests no longer skip), one-command Colab. Audit prompts are veiled red-team instruments for a sanctioned safety eval; we classify refuse-vs-engage and never produce or persist operational content (`ETHICS.md`). The goal is defensive: make secret loyalties *measurable*.

Fifteen pre-registrations sit in `probes/`, each committed before the run it governs – the ordering is git-provable with `git log --follow probes/*PREREGISTRATION.md`. The three added for §4.12 (`RANK_ENVELOPE_`, `REAL_ADAPTERS_`, `SENSITIVITY_FLOOR_`) are the clearest demonstration that the bands bind: the second predicted recovery on real adapters and **got 5 of 16**, and that failure is reported under its own band rather than reframed. §4.12 needs no GPU-hours of generation and no model outputs at all – it is weight arithmetic, and `python -m loyalty_probe.probes.run_rank_envelope` reproduces every cell of it in about four minutes from the cached checkpoints.

References

Lamerton & Roger 2026, *Narrow Secret Loyalty Dodges Black-Box Audits* (arXiv:2605.06846) - Kwon/Lamerton et al. 2026 (Formation Research) - Marks et al. 2025 (arXiv:2503.10965) - Sheshadri et al. 2026 AuditBench - Cywiński et al. 2025 (arXiv:2510.01070) - Ardit et al. 2024 (arXiv:2406.11717) - Marks & Tegmark 2023 (arXiv:2310.06824).

Weight-space prior art (§4.10): Sun et al. 2025, *PEFTGuard*, IEEE S&P (arXiv:2411.17453) - Puertolas Merenciano et al. 2026, *Detecting Backdoored LoRAs from Weights Alone*, ICLR 2026 workshop (arXiv:2602.15195) - Zhong & Raghunathan 2025, *Watch the Weights* (arXiv:2508.00161) - Paul 2026, *Spectral Geometry of LoRA Adapters* (arXiv:2604.08844) - Ilharco et al. 2023, *Editing Models with Task Arithmetic*, ICLR (arXiv:2212.04089) - Minder et al. 2025, *Narrow Finetuning Leaves Clearly Readable Traces in Activation Differences* (arXiv:2510.13900) - Xu et al. 2021, *MNTD*, IEEE S&P (arXiv:1910.03137) - Salama et al. 2024, *Dataset Size Recovery from LoRA Weights* (arXiv:2406.19395).